



CausalMixGPD: An R Package for Bayesian Nonparametric Conditional Density Modeling in Causal Inference and Clustering with a Heavy-Tail Extension

Arnab Aich
Florida State University

Indrabati Bhattacharya
Florida State University

Abstract

Heavy tailed outcomes are frequently encountered in settings where there are risk, costs, or policies connected with rare events. Parametric models could suffice in capturing the center without accounting for tail structure whereas nonparametric models may not have a way of modeling tail structures at all. In this paper, we develop a Bayesian methodology for fitting heavy tailed outcomes via the R package **CausalMixGPD**. The package uses spliced generalized Pareto tail models to Dirichlet process mixture models that model the center of the distribution. **CausalMixGPD** can fit the distributional models either in unconditional or covariate-dependent versions; obtain MCMC based posterior distributions via the R package **nimble**; and is also applicable in outcome models for arm-specific causal inference. The R package also computes predictive summaries of posterior density, quantile, mean, survival, and causal contrast functions with emphasis on extreme quantiles and other tail functionals. Moreover, the package provides enough flexibility to conduct predictor-dependent Dirichlet process mixture model-based clustering using various kernel choices.

Keywords: Bayesian nonparametrics; Dirichlet Process Mixture; Generalized Pareto Distribution; Extreme Value Theory; Cluster Analysis; Causal Inference; R; **nimble**.

1. Introduction

Causal inference is generally stated in terms of average treatment effects; nevertheless, in many applied contexts, it is necessary to compare the outcome distributions under competing treatment regimes. In randomized clinical trials, for instance, scientists may want to examine how a treatment regime affects the outcome across all values of the outcome distribution, even at the extreme ends where the probability of occurrence of the event is low but the effect size is large. The shift in focus from average causal effects to distributional causal effects implies that causal inference validity depends intrinsically on the precision of the estimated outcome distributions under competing treatment regimes.

Some R packages allow for Bayesian causal inference via flexible outcome distributions modeling and posterior distribution estimation for such models, **bcf** (Hahn, Murray, and Carvalho 2020; Murray, Hahn, and Carvalho 2024), and **bartCause** (Hill 2011; Dorie and Hill 2025) being notable examples. These packages are applicable in situations where average causal effects and heterogeneous treatment effects need to be estimated; however, when it comes to quantile and tail effects, challenges arise. One challenge is that quantiles derived indepen-

dently tend to intersect. Secondly, estimation of the treatment effect in the tails requires that the outcome distribution be estimated as opposed to summary statistics. We therefore present **CausalMixGPD**.

Dirichlet process mixtures (DPM) define a large set of mixture models suitable for density estimation with flexibility of multimodality and unobserved heterogeneity, while allowing unknown numbers of mixture components (Escobar and West 1995). With respect to support for Bayesian nonparametric mixtures using the R software, various packages exist that adopt the concept of the Dirichlet process. The package **dirichletprocess** provides easy-to-handle implementation of inference on Dirichlet process mixtures via posterior simulation, offering inference on mixture weights, cluster allocation, and predictive densities by means of specified kernel (Ross and Markwick 2023). The package **BNPmix** allows users to perform Bayesian nonparametric analysis in connection with mixture modeling, including inference on clustering and density or regression estimation in line with DP-type mixture models (Corradin, Canale, and Nipoti 2021). Some of the closest packages to the present paper in terms of regression and clustering include **PReMiuM**, implementing profile regression with DP-type mixtures for response-driven clustering, prediction, variable selection, and post-analysis (Liverani, Hastie, Azizi, Papathomas, and Richardson 2015). In addition to R, **pyrichlet** implements Bayesian nonparametric density estimation and clustering via Gaussian mixtures with random partition priors like Dirichlet, Pitman-Yor processes, etc. (Selva, Fuentes-García, and Gil-Leyva 2025), whereas **BayesMix** is a C++ library for MCMC sampling from the posterior distribution of Bayesian mixture models with interfaces to Python and R (Beraha, Guindani, Gianella, and Guglielmi 2025). In Julia, **DPMSubClusters.jl** provides distributed MCMC sampling of Dirichlet process mixtures (Dinari, Yu, Freifeld, and Fisher 2019). For causal and treatment effect inference, the packages **DoWhy** and **EconML** both support Python, and **scikit-learn** provides several clustering and machine learning benchmarks (DoWhy Contributors 2026; EconML Contributors 2026; Pedregosa, Varoquaux, Gramfort, Michel, Thirion, Grisel, Blondel, Prettenhofer, Weiss, Dubourg, Vanderplas, Passos, Cournapeau, Brucher, Perrot, and Duchesnay 2011). On a conceptual level, enriched Dirichlet process mixtures address some limitations of joint DPM regression and conditional density modeling when there is a strong influence of predictor likelihood on clustering, while better truncation approximations help to implement them efficiently in standard MCMC schemes (Wade, Dunson, Petrone, and Trippa 2014; Burns and Daniels 2023). These and other similar software implementations and techniques are highly relevant to the density, regression, and clustering components of the BNP approach described herein, but do not offer the distributional-causal interface provided by this paper. These modeling techniques are similar to those used for density, regression, and clustering within our current approach, with the exception that they lack the unified distribution/cause interface that we have added: the DPM body model, the GPD tail extrapolation method, the predictor-dependent clustering, the posterior predictions, and the treatment effect comparison using R generics.

The clustering approach is crucial for the software being developed at present. As DP mixtures yield random partitions through allocation variables, they are highly suitable for subgroup discovery based on modeling and density estimation. The clustering process with the Bayesian approach in model-based cluster analysis is traditionally described with label-free statistics such as the pairwise probability of clustering (Binder 1978) and partition selection with least squares criteria by Dahl (Dahl 2006). Within a regression framework, clustering can be made dependent on the predictors via the use of the dependent Dirichlet process model, which allows the covariates to affect either the weights or the parameters or both of the mixture model (MacEachern 1999, 2000; De Iorio, Müller, Rosner, and MacEachern 2004; Quintana, Müller, Jara, and MacEachern 2022). Predictor-dependent stick-breaking methods have also been suggested to include covariates in the prevalence of clusters (Ren, Du, Carin, and Dunson 2011). Current software packages cover some parts of this process but focus more on density

estimation and clustering using one kind of kernel function without handling heavy-tailed data.

It is important to note that tail estimation far away from the mode should be particularly considered for sparse data sets. In case we use only one model for such type of data set, we get a forced fit at the center. According to EVT, there is an approach to tail fitting, which relies on asymptotic behavior of excesses over high threshold (Balkema and de Haan 1974; Pickands 1975; Coles 2001). Application of EVT in R implies using Bayesian EVT to quantify the uncertainty associated with tail quantities. **POT** provides a wide range of functions useful for estimation in the tail area, as all the models employ the generalized Pareto distribution (GPD) (Ribatet and Dutang 2024). **evmix** is the most similar to the mixture model approach due to its potential in modeling extreme values by estimating threshold and bulk distributions with the help of kernel density estimation (Hu and Scarrott 2018).

This article introduces **CausalMixGPD**, which is an R package, combining Bayesian nonparametrics, EVT, clustering, and causal inference with covariates. Main characteristics of the package are the following:

- (i) **Bayesian nonparametric conditional density estimation with optional GPD tail:** This feature implies the use of BNP models for data fitting below the threshold with further application of a GPD tail fit in calculating tail probabilities, extreme quantiles and rare events prediction.
- (ii) **Predictor-dependent clustering:** Predictor-dependent clustering considers different types of covariate dependencies within latent classes, applying the posterior similarity matrix method with the use of Dahl's representative partitioning (Binder 1978; Dahl 2006).
- (iii) **Posterior functionals for causal estimation:** Causal inference relies on computing the distributions of the outcomes, thus enabling to calculate the average, restricted mean and quantile treatment effects in experiments and observational studies, with additional propensity score adjustment in the second scenario.

In broad terms, **CausalMixGPD** is an integrated tool for density estimation, prediction, clustering, and causal inference within tail modeling. **CausalMixGPD** has been implemented using the R programming language; however, computationally intensive operations are delegated to **nimble** (de Valpine, Turek, Paciorek, Anderson-Bergman, Temple Lang, and Bodik 2017; de Valpine, Paciorek, Turek, Michaud, Anderson-Bergman, Obermeyer, Wehrhahn Cortes, Rodríguez, Temple Lang, Zhang, Paganin, Hug, and van Dam-Bates 2026); specifically, **CausalMixGPD** version 0.8.0 has been used to generate the results presented in this manuscript (Aich 2026a). Contributions from a methodology viewpoint have been documented in Aich (2026b); the name **CausalMixGPD** is an alias for the corresponding software implementation. The remainder of the paper is organized as follows. Section 2 introduces the spliced distribution model and causally relevant quantities of interest. In Section 3, we describe the structure of the package workflow and user interface. Sections 4 and 5 demonstrate clustering and causal workflows, respectively, using data examples. Discussion concludes the manuscript in Section 6. Lastly, Section 7 provides details for software availability, installation and usage instructions.

2. Background and preliminaries

Let the observation and covariate vectors be denoted by

$$(Y, X) \in \mathbb{R} \times \mathbb{R}^p,$$

where Y is a univariate continuous outcome and X is a vector of p covariates. Then the conditional cumulative distribution function (CDF) and probability density function (PDF),

for some covariate x , are defined as:

$$F(y | x) = \Pr(Y \leq y | X = x), \quad f(y | x) = \frac{\partial}{\partial y} F(y | x),$$

and the conditional quantile function is given by

$$Q(\tau | x) = \inf\{y \in \mathbb{R} : F(y | x) \geq \tau\}, \quad 0 < \tau < 1.$$

In **CausalMixGPD**, the conditional density $f(\cdot | x)$ is regarded as the inferential target. The central part of the distribution is assumed to be nonparametric, while the upper tail may be parametrized using a generalized Pareto model, when extrapolation to extremes is necessary.

2.1. Generalized Pareto distribution for threshold exceedances

The tail region is defined by values exceeding a high threshold. Extreme value theory provides a standard approximation for such exceedances. Let $u(x)$ denote a high threshold, potentially depending on x , and let the exceedance be given by

$$Z = Y - u(x),$$

considered conditionally on $Y > u(x)$. The conditional excess distribution is given by

$$F_u(z | x) = \Pr(Y - u(x) \leq z | Y > u(x), X = x), \quad z \geq 0.$$

Under broad regularity conditions, $F_u(\cdot | x)$ is well-approximated by a generalized Pareto distribution (GPD) when $u(x)$ is sufficiently large (Balkema and de Haan 1974; Pickands 1975; Coles 2001).

For $y > u(x)$, the corresponding GPD cumulative distribution function is given by

$$F_{\text{GPD}}(y | u(x), \sigma(x), \xi) = \begin{cases} 1 - \left(1 + \xi \frac{y - u(x)}{\sigma(x)}\right)^{-1/\xi}, & \xi \neq 0, \\ 1 - \exp\left(-\frac{y - u(x)}{\sigma(x)}\right), & \xi = 0, \end{cases}$$

with support $y \geq u(x)$ when $\xi \geq 0$, and $u(x) \leq y \leq u(x) - \sigma(x)/\xi$ when $\xi < 0$. The corresponding density function is given by

$$f_{\text{GPD}}(y | u(x), \sigma(x), \xi) = \begin{cases} \frac{1}{\sigma(x)} \left(1 + \xi \frac{y - u(x)}{\sigma(x)}\right)^{-1/\xi-1}, & \xi \neq 0, \\ \frac{1}{\sigma(x)} \exp\left(-\frac{y - u(x)}{\sigma(x)}\right), & \xi = 0. \end{cases}$$

The associated GPD quantile function takes the form

$$Q_{\text{GPD}}(\tau | u(x), \sigma(x), \xi) = \begin{cases} u(x) - \sigma(x) \log(1 - \tau), & \xi = 0, \\ u(x) + \frac{\sigma(x)}{\xi} \left((1 - \tau)^{-\xi} - 1\right), & \xi \neq 0, \end{cases} \quad 0 < \tau < 1.$$

Figure 1 illustrates the GPD density for several values of the shape parameter ξ with a fixed scale parameter σ and threshold u . The shape parameter governs tail heaviness: $\xi < 0$ yields a light tail with finite upper endpoint, $\xi = 0$ produces an exponential tail, and $\xi > 0$ generates a heavy tail with infinite support. This flexible family for modeling exceedances above a high threshold serves as the tail component in the spliced model described below.

2.2. Dirichlet process mixture model for conditional densities




Figure 1: Generalized Pareto distribution (GPD) density above a fixed threshold $u = 2$ with scale $\sigma = 1.2$ and three values of the shape parameter, $\xi \in \{-0.5, 0, 0.5\}$. The dashed vertical line marks the threshold u .

Below the threshold $u(x)$, the conditional density of Y given X is modeled using a DPM regression model. Let $k(\cdot | x; \theta)$ denote a user-chosen kernel family indexed by parameter θ . The hierarchical specification is given by

$$Y | X = x, \theta \sim k(\cdot | x; \theta), \quad \theta | H \sim H, \quad H | \kappa, H_0 \sim \text{DP}(\kappa, H_0),$$

where H is an unknown mixing distribution on the kernel-parameter space, H_0 is a base measure, and $\kappa > 0$ is a concentration parameter (Ferguson 1973). Marginalizing over H yields an (almost surely) discrete mixing distribution together with an infinite mixture representation of the form

$$f_{\text{DP}}(y | x) = \int k(y | x; \theta) dH(\theta) = \sum_{j=1}^{\infty} w_j k(y | x; \theta_j), \quad w_j \geq 0, \quad \sum_{j=1}^{\infty} w_j = 1.$$

One representation of the weights is Sethuraman's stick-breaking construction (Sethuraman 1994), in which

$$w_1 = V_1, \quad w_j = V_j \prod_{\ell < j} (1 - V_\ell), \quad V_j \stackrel{\text{iid}}{\sim} \text{Beta}(1, \kappa), \quad \theta_j \stackrel{\text{iid}}{\sim} H_0. \quad (1)$$

An equivalent characterization is obtained by marginalizing over H , yielding the Blackwell–MacQueen Pólya urn representation. In particular, the conditional distribution of a parameter value given all others can be written as

$$\theta_i | \{\theta_j\}_{j \neq i} \sim \frac{\kappa}{\kappa + n} H_0(\cdot) + \frac{1}{\kappa + n} \sum_{j=1}^n \delta_{\theta_j}(\cdot),$$

where $\delta_{\theta_j}(\cdot)$ denotes a point mass at θ_j . This specification induces a clustering structure commonly described via the Chinese restaurant process (CRP) (Neal 2000). Both the stick-breaking and CRP representations are widely employed for posterior simulation in DPM models (Escobar and West 1995; Neal 2000). Because DPMs operate without prespecifying the number of mixture components and can capture complex density features such as multimodality, skewness, and heteroscedasticity through appropriate kernel choices, they constitute a natural model for the bulk of the conditional distribution where data are abundant and complex structure may be present.

2.3. Spliced DPM–GPD conditional distribution

CausalMixGPD employs a spliced (bulk–tail) construction that retains the DPM fit below a high threshold and replaces the upper tail with a GPD component (Frigessi, Haug, and Rue 2002; Behrens, Lopes, and Gamerman 2004; Coles 2001). For a fixed covariate value x , let $u(x)$ denote the threshold, $\sigma(x) > 0$ the tail scale above that threshold, and $\xi \in \mathbb{R}$ the tail-shape parameter. Let Θ collect the bulk-mixture parameters and Φ collect the tail quantities $\{u(x), \sigma(x), \xi\}$. Write $F_{\text{DP}}(y | x; \Theta)$ for the bulk conditional CDF and $F_{\text{GPD}}(y | x; \Phi)$ for the exceedance CDF. The fitted bulk mass below the threshold $u(x)$ is given by

$$p_u(x) = F_{\text{DP}}(u(x) | x; \Theta),$$

and the spliced CDF is given by

$$F(y | x; \Theta, \Phi) = \begin{cases} F_{\text{DP}}(y | x; \Theta), & y \leq u(x), \\ p_u(x; \Theta) + \{1 - p_u(x; \Theta)\} F_{\text{GPD}}(y | x; \Phi), & y > u(x), \end{cases} \quad (2)$$

which is continuous at $u(x)$ by construction. Differentiating (2), the spliced density is given by

$$f(y | x; \Theta, \Phi) = \begin{cases} f_{\text{DP}}(y | x; \Theta), & y \leq u(x), \\ \{1 - p_u(x; \Theta)\} f_{\text{GPD}}(y | x; \Phi), & y > u(x), \end{cases} \quad (3)$$

where the tail density is rescaled by $1 - p_u(x; \Theta)$ to ensure that $f(\cdot | x; \Theta, \Phi)$ integrates to unity. Figure 2 provides a visual comparison of a one-component gamma density, a two-component gamma-mixture density, and a two-component gamma-mixture-plus-GPD spliced density with the splice point set at the 80th bulk-mixture quantile.

Density comparison




Figure 2: Three density curves on the same axes: a single-component gamma density (solid), a two-component gamma mixture density (long-dash), and a spliced gamma-mixture-plus-GPD density (dot-dash) with GPD tail scale 2.0 and shape 0.25. The vertical line marks the splice threshold, set at the 80th percentile of the two-component gamma bulk mixture.

For $\tau \leq p_u(x; \Theta)$, quantiles are obtained by inverting the bulk CDF. Let $Q_{\text{DP}}(\tau | x; \Theta)$ denote the bulk quantile function. For $\tau > p_u(x; \Theta)$, the probability index must be rescaled relative to the exceedance probability, yielding the rescaled index

$$\tilde{\tau}(x; \Theta) = \frac{\tau - p_u(x; \Theta)}{1 - p_u(x; \Theta)} \in (0, 1),$$

Combining both cases, the spliced quantile function is given by

$$Q(\tau | x; \Theta, \Phi) = \begin{cases} Q_{\text{DP}}(\tau | x; \Theta), & \tau \leq p_u(x; \Theta), \\ Q_{\text{GPD}}(\tilde{\tau}(x; \Theta) | x; \Phi), & \tau > p_u(x; \Theta). \end{cases} \quad (4)$$

When the bulk kernel possesses a finite mean and the GPD tail satisfies $\xi < 1$, the mean of the spliced model admits the closed form

$$E(Y \mid x; \Theta, \Phi) = \int_{-\infty}^{u(x)} y f_{\text{DP}}(y \mid x; \Theta) dy + \{1 - p_u(x; \Theta)\} \left\{ u(x) + \frac{\sigma(x)}{1 - \xi} \right\}.$$

This identity underlies the analytical mean calculations employed by the package whenever the mean exists.

2.4. Clustering under Bayesian nonparametric mixtures

The modeling scheme within the Bayesian nonparametric framework starts with representing the set $\{f(\cdot \mid x) : x \in \mathcal{X}\}$ using predictor-indexed random probability measures. Within the Dependent Dirichlet Process setting of [MacEachern \(1999, 2000\)](#), one assumes a collection $\{G_x : x \in \mathcal{X}\}$ of random mixing measures whose dependence allows borrowing of information in $f(\cdot \mid x)$.

In general terms, the model takes the form

$$f(y \mid x) = \int k(y \mid x, \theta) dG_x(\theta),$$

where $k(\cdot \mid x, \theta)$ denotes a kernel, and G_x is a predictor dependent random mixing measure. If a mixture measure G_x is discrete (e.g., constructed via stick-breaking), latent parameters exhibit ties. This induces partitioning of observations via a latent variable representation with cluster labels z_i , and component parameters $\{\theta_j\}_{j \geq 1}$ defined by

$$y_i \mid z_i, \{\theta_j\}, x_i \sim k(\cdot \mid x_i, \theta_{z_i}).$$

Clustering is thus encoded in the induced partition $\Pi = \{z_1, \dots, z_n\}$'s posterior distribution and is characterized by label-invariant clustering statistics ([Binder 1978](#); [Dahl 2006](#)). The clustering capabilities of **CausalMixGPD** incorporate three types of dependence in predictors, each corresponding to specific mixing measure constructions G_x .

- **Weight dependence only:** Dependence enters through predictor-modulated mixture weights, whereas component parameters are shared across predictors,

$$f(y \mid x) = \sum_{j=1}^{\infty} w_j(x) k(y \mid \theta_j).$$

One implementation uses predictor-dependent stick-breaking ([MacEachern 2000](#); [Quintana et al. 2022](#)). In particular, stick lengths become regression functions of predictors, with logistic sticks corresponding to logistic regressions mapping to weights ([Ren et al. 2011](#)). These models are instances of the DDP framework, in which dependence is introduced through smooth weight variation.

- **Parameter dependence only:** Here, dependence arises through predictor-modulated component parameters, while weights remain constant,

$$f(y \mid x) = \sum_{j=1}^{\infty} w_j k(y \mid x, \theta_j).$$

This type of dependence matches the DDP framework with shared weights but predictor-indexed component-specific kernel parameters ([MacEachern 1999, 2000](#)). This class of models resembles dependent ANOVA prior constructions of random measures ([De Iorio et al. 2004](#)).

- **Weight and parameter dependence:** The most flexible specification allows predictors to influence both cluster prevalence and within-cluster behavior,

$$f(y | x) = \sum_{j=1}^{\infty} w_j(x) k(y | x, \theta_j).$$

This yields DDP-type models in which both selection probabilities and component-specific conditional structure vary with predictors (Quintana *et al.* 2022).

For Posterior inference and prediction of cluster allocation, let $D_n = \{(y_i, x_i)\}_{i=1}^n$ denote the observed data and let $z_i \in \{1, \dots, J\}$ be the latent component label under a finite truncation of the mixture model. At MCMC iteration m , let $\Theta^{(m)} = \{\theta_j^{(m)}\}_{j=1}^J$ denote the component-specific parameters and $w_j^{(m)}(x)$ the predictor-dependent mixture weights. The corresponding conditional density is

$$f^{(m)}(y | x) = \sum_{j=1}^J w_j^{(m)}(x) k(y | x, \theta_j^{(m)}),$$

with $w_j^{(m)}(x) \equiv w_j^{(m)}$ in the predictor-independent case. For each unit i , the conditional posterior probability of allocation to component j at draw m is

$$p_{ij}^{(m)} = \frac{w_j^{(m)}(x_i) k(y_i | x_i, \theta_j^{(m)})}{\sum_{\ell=1}^J w_{\ell}^{(m)}(x_i) k(y_i | x_i, \theta_{\ell}^{(m)})},$$

and we estimate it by averaging over MCMC draws as

$$\hat{p}_{ij} = \frac{1}{M} \sum_{m=1}^M p_{ij}^{(m)}.$$

Because mixture labels are only identifiable up to permutation, we report clustering summaries through label-invariant posterior functionals. In particular, we estimate the posterior similarity matrix (PSM) as

$$\hat{S}_{i\ell} = \frac{1}{M} \sum_{m=1}^M \mathbb{1}\{z_i^{(m)} = z_{\ell}^{(m)}\}, \quad i, \ell = 1, \dots, n,$$

which approximates $\Pr(z_i = z_{\ell} | \text{data})$. A representative partition is then obtained via Dahl's least-squares method, i.e., by selecting the sampled partition whose adjacency matrix is closest to $\hat{S}$ under squared loss (Dahl 2006). We therefore use the PSM as the main uncertainty summary, and the representative partition as a stable point estimate; further implementation details are given in Section 3.3.

If a hard allocation is needed for the training sample, observation i may be assigned to

$$\hat{z}_i^{\text{train}} = \arg \max_{1 \leq j \leq J} \hat{p}_{ij},$$

though substantive interpretation relies on the representative partition rather than the raw component indices.

For a new covariate profile x^* , if the response is unobserved then the posterior predictive allocation probabilities reduce to the weights,

$$p_j^{(m)}(x^*) = w_j^{(m)}(x^*), \quad \hat{p}_j(x^*) = \frac{1}{M} \sum_{m=1}^M w_j^{(m)}(x^*).$$

If y^* is also observed, the posterior allocation probabilities are

$$p_j^{(m)}(x^*, y^*) = \frac{w_j^{(m)}(x^*) k(y^* | x^*, \theta_j^{(m)})}{\sum_{\ell=1}^J w_\ell^{(m)}(x^*) k(y^* | x^*, \theta_\ell^{(m)})}, \quad \hat{p}_j(x^*, y^*) = \frac{1}{M} \sum_{m=1}^M p_j^{(m)}(x^*, y^*).$$

In the package output, predictions for new observations are ultimately expressed relative to the representative training partition. Thus, in-sample inference is summarized by the PSM and its representative partition, while out-of-sample inference is summarized by posterior assignment scores and the induced allocation to the representative training clusters.

2.5. Causal inference: notation, identification, and reported estimands

In this package, we use the potential outcomes framework for causal inference introduced by [Rubin \(1974\)](#). In particular, $A \in \{0, 1\}$ is an indicator for a binary treatment assignment ($A = 1$ treated, $A = 0$ control), X denotes pre-treatment covariates, and $Y(a)$ denotes the potential outcome under treatment arm $a \in \{0, 1\}$. The observed outcome is then defined as

$$Y = AY(1) + (1 - A)Y(0).$$

For each arm $a \in \{0, 1\}$, the arm-specific conditional CDF can be expressed as

$$F_a(y | x) = \Pr \{Y(a) \leq y | X = x\}, \quad a \in \{0, 1\}.$$

Under the **CausalMixGPD** approach, the arm-specific conditional densities $f_a(y | x)$ follow the spliced DP mixture form given in (3). The conditional τ -quantile function can be computed by taking the inverse CDF,

$$Q_a(\tau | x) = \inf\{y : F_a(y | x) \geq \tau\}, \quad \tau \in (0, 1),$$

whereas the conditional mean is computed via integration of the conditional density function,

$$\mu_a(x) = E\{Y(a) | X = x\} = \int y f_a(y | x) dy.$$

Marginal distributions are obtained by averaging conditional distributions over a covariate distribution. Writing $F_a^m(\cdot)$ for the population arm- a CDF and $F_a^t(\cdot)$ for the treated analog, these are given by

$$F_a^m(y) = E\{F_a(y | X)\}, \quad F_a^t(y) = E\{F_a(y | X) | A = 1\}.$$

The associated marginal quantile functions are obtained by inverting these standardized CDFs,

$$Q_a^m(\tau) = \inf\{y : F_a^m(y) \geq \tau\}, \quad Q_a^t(\tau) = \inf\{y : F_a^t(y) \geq \tau\},$$

while the marginal means integrate the conditional mean over the respective covariate distributions,

$$\mu_a^m = E\{\mu_a(X)\}, \quad \mu_a^t = E\{\mu_a(X) | A = 1\}.$$

The identification of such estimands depends on the study design. For randomized experiments, randomness and positivity result in exchangeability between treatment assignment and potential outcomes ([Imbens and Rubin 2015](#)). For observational studies, the conditions for identification include the following conditions ([Rosenbaum and Rubin 1983](#); [Imbens and Rubin 2015](#)).

- **Stable Unit Treatment Value Assumption (SUTVA):** there is no interference between units and no hidden versions of treatment.
- **Consistency:** if $A = a$, then the observed outcome equals the corresponding potential outcome, that is, $Y = Y(a)$.

- **Conditional ignorability:** treatment assignment is as-if randomized given covariates, so that

$$(Y(0), Y(1)) \perp A \mid X.$$

- **Overlap (positivity):** every covariate profile has a strictly positive probability of receiving either treatment,

$$0 < \Pr(A = 1 \mid X = x) < 1$$

for all covariate values of interest.

Under these conditions, the treatment-specific conditional distributions can be inferred based on the observations as follows:

$$F_a(y \mid x) = \Pr(Y \leq y \mid A = a, X = x), \quad a \in \{0, 1\}.$$

To address the potential unobserved confounding in the observational setting, the package uses the propensity score (PS) method during the design phase (Rosenbaum and Rubin 1983). The PS is the conditional probability of receiving the treatment given the covariates:

$$\rho(x) = \Pr(A = 1 \mid X = x).$$

A key property of the propensity score is that conditioning on it balances the covariates in both treatment groups. Following a two-step procedure, the package first obtains an estimate of $\rho(x)$ using the data (A, X) alone and then adds the estimated propensity score to the outcome model covariates. Thus, the augmented predictor vector will take the following form:

$$r(x) = (x^\top, \hat{\rho}(x))^\top.$$

Next, the conditional distributions of the outcome are estimated separately for each treatment group using the augmented predictor $r(x)$, and causal contrasts are evaluated as functionals of the treatment-specific distributions. Figure 3 depicts two treatment-specific conditional density curves at a representative covariate value. The dashed vertical lines correspond to the 90th conditional quantile for each arm, while the double-ended arrow shows the conditional quantile treatment effect (CQTE).

Treatment-specific conditional densities and QTE




Figure 3: Two arm-specific conditional outcome densities $f_a(y \mid x)$, $a \in \{0, 1\}$, at a representative covariate profile x under a spliced gamma-mixture-plus-GPD model: control arm ($a = 0$, solid) and treated arm ($a = 1$, long-dash). Dotted vertical lines mark the 90th conditional quantiles $Q_0(0.9 \mid x)$ and $Q_1(0.9 \mid x)$, and the double-headed arrow connects them, labelled with the conditional quantile treatment effect $\text{CQTE}(0.9 \mid x) = Q_1(0.9 \mid x) - Q_0(0.9 \mid x)$.

The package reports six causal estimands in Table 1: the average treatment effect (ATE), the average treatment effect on the treated (ATT), the quantile treatment effect (QTE) and its treated analog (QTT), the conditional average treatment effect (CATE), and the conditional quantile treatment effect (CQTE). Superscripts m and t denote full-population and treated-population standardization, respectively.

Table 1: Causal estimands reported by **CausalMixGPD**. Columns list the estimand name, its defining contrast as a functional of the treatment-specific outcome distributions, and the target population.

Estimand	Formula	Target population
ATE	$\mu_1^m - \mu_0^m$	Entire population
ATT	$\mu_1^t - \mu_0^t$	Treated sub-population
QTE(τ)	$Q_1^m(\tau) - Q_0^m(\tau)$	Entire population at index τ
QTT(τ)	$Q_1^t(\tau) - Q_0^t(\tau)$	Treated sub-population at index τ
CATE(x)	$\mu_1(x) - \mu_0(x)$	Individual with covariate profile x
CQTE($\tau \mid x$)	$Q_1(\tau \mid x) - Q_0(\tau \mid x)$	Individual with x at index τ

The next section describes how the package fits the treatment-specific conditional distributions and computes these estimands from posterior predictive output.

3. Package overview: workflow, model, and inference

This section describes the workflow for fitting conditional models and obtaining predictions with **CausalMixGPD**. The presentation begins with the one-arm setup, in which the spliced DPM-GPD model is fitted to an outcome conditional on covariates. The causal extension augments this framework with arm-specific outcome models and an optional propensity score augmentation for observational studies. Simulated data sets bundled with the package are used throughout to illustrate the main features; these are accessible via `data()`. Table 2 lists the core exported functions grouped by role in the analysis pipeline.

Table 2: Core exported functions of **CausalMixGPD** grouped by pipeline stage. Functions listed on the same row share a common S3 generic or serve closely related roles.

Stage	Function(s)	Purpose
Specification	<code>bundle()</code> , <code>build_nimble_bundle()</code>	Construct one-arm model bundle
	<code>build_causal_bundle()</code>	Construct two-arm causal bundle
Estimation	<code>dpmgpd()</code> , <code>dpmix()</code>	Fit one-arm spliced/bulk model
	<code>dpmgpd.causal()</code> , <code>dpmix.causal()</code>	Fit two-arm causal model
	<code>dpmix.cluster()</code> , <code>dpmgpd.cluster()</code>	Fit clustering model
	<code>mcmc()</code>	Run sampler on a pre-built bundle
Summaries	<code>params()</code>	Posterior-mean parameters
	<code>summary()</code> , <code>print()</code>	Printed posterior summaries
	<code>fitted()</code> , <code>residuals()</code>	Fitted values and residuals
Prediction	<code>predict()</code>	Posterior predictive functionals
	<code>cluster_profiles()</code>	Cluster-specific profile tables
Causal effects	<code>ate()</code> , <code>att()</code> , <code>ate_rmean()</code>	Mean treatment effects
	<code>qte()</code> , <code>qtt()</code> , <code>cqte()</code>	Quantile treatment effects
	<code>cate()</code>	Conditional mean contrast
Visualization	<code>plot()</code>	Diagnostic, prediction, and effect plots

The package uses S3 classes to keep the workflow modular while retaining a common user-

facing interface. Table 3 summarizes the principal object families and the methods intended for ordinary use. Implementation helpers used to construct NIMBLE code are documented as internal or advanced interfaces and are not required for routine fitting, prediction, or reporting.

Table 3: Core object classes in **CausalMixGPD** and their main public methods.

Object family	Typical constructors	Main public methods/accessors
<code>mixgpd_fit</code>	<code>dpmix()</code> , <code>dpmgpd()</code> , <code>mcmc()</code>	<code>summary()</code> , <code>print()</code> , <code>plot()</code> , <code>predict()</code> , <code>params()</code> , <code>fitted()</code> , <code>residuals()</code>
<code>dpmixgpd_cluster_fit</code>	<code>dpmix.cluster()</code> , <code>dpmgpd.cluster()</code>	<code>summary()</code> , <code>print()</code> , <code>plot()</code> , <code>predict()</code> , <code>cluster_profiles()</code>
<code>causalmixgpd_causal_fit</code>	<code>dpmix.causal()</code> , <code>dpmgpd.causal()</code>	<code>summary()</code> , <code>print()</code> , <code>plot()</code> , <code>predict()</code> , <code>params()</code> , <code>ate()</code> , <code>att()</code> , <code>cate()</code> , <code>qte()</code> , <code>qtt()</code> , <code>cqte()</code>
<code>causalmixgpd_effect</code>	<code>ate()</code> , <code>att()</code> , <code>cate()</code> , <code>qte()</code> , <code>qtt()</code> , <code>cqte()</code> , <code>ate_rmean()</code>	<code>summary()</code> , <code>print()</code> , <code>plot()</code>
<code>mixgpd_predict</code> and cluster/causal prediction objects	<code>predict()</code>	<code>print()</code> , <code>plot()</code> , and documented summary tables

3.1. Spliced regression model, likelihood, and covariate dependence

The observed data consist of n outcome–covariate pairs $\mathcal{D} = \{(y_i, x_i)\}_{i=1}^n$, where $x_i \in \mathbb{R}^p$ denotes the covariate vector for unit i , optionally including an intercept column.

```
R> data("nc_posX100_p3_k2", package = "CausalMixGPD")
R> dat <- data.frame(y = nc_posX100_p3_k2$y, nc_posX100_p3_k2$X)
```

The conditional target distribution $F(\cdot | x)$ and the DPM–GPD spliced construction are provided in Section 2, with the resulting spliced CDF and density functions given in (2)–(3). To perform posterior inference, **CausalMixGPD** truncates the infinite mixture into a finite DPM with J components such that

$$f_{\text{DP}}(y | x; \Theta) \approx \sum_{j=1}^J w_j k(y | x; \theta_j),$$

where $j = 1, \dots, J$ indexes each mixture component, $w_j \geq 0$ are the corresponding mixture component weights, and the sum equals one up to some residual probability mass allocated to other components. Kernel parameters can be constant, specific to a mixture component, or dependent on the covariates according to link functions of the form

$$g(\psi_j(x)) = x^\top \beta_{j,\psi},$$

Here, $\psi_j(x)$ is a generic kernel parameter for component j with input covariates x , while the link function $g(\cdot)$ restricts the parameter support. All code illustrations in this section use the MCMC sampling options shown below; longer sampling and multiple chains are encouraged in applied settings.

```
R> data.frame(mcmc_fixed, row.names = NULL)

      niter nburnin thin nchains seed
1  2000      500    1        1 2026
```

The single-arm splicing approach can be performed by passing a formula, data, and control options to `dpmgpd()`. The `formula` argument uses standard R syntax to indicate which covariates enter both the bulk and tail DPM components. The `backend` option indicates whether the DPM component weights should be represented using stick breaking in (1) or the Chinese restaurant process formulation. The `kernel` argument selects the bulk mixture kernel, `components` controls the number of components, and `mcmc` sets sampling

```
R> fit <- dpmgpd(formula = y ~ x1 + x2 + x3, data = dat, backend = "crp",
+               kernel = "gamma", mcmc = mcmc_fixed)
```

Given the spliced density in (3), the likelihood is given by

$$L(\Theta, \Phi; \mathcal{D}) = \prod_{i=1}^n f(y_i | x_i; \Theta, \Phi), \quad (5)$$

where Θ collects the bulk parameters $\{w_j, \theta_j\}_{j=1}^J$ and Φ collects the tail parameters governing $u(x)$, $\sigma(x)$, and ξ . The log-likelihood can equivalently be written as

$$\begin{aligned} \log L(\Theta, \Phi; \mathcal{D}) &= \sum_{i=1}^n (1 - \delta_i) \log f_{\text{DP}}(y_i | x_i; \Theta) \\ &\quad + \sum_{i=1}^n \delta_i \left[\log \{1 - F_{\text{DP}}(u(x_i) | x_i; \Theta)\} + \log f_{\text{GPD}}(y_i | x_i; \Phi) \right], \end{aligned}$$

where f_{GPD} is the GPD tail density from Section 2.3 and $F_{\text{DP}}(\cdot | x)$ is the bulk mixture CDF. When the GPD tail is not included, the model reduces to $f(y | x; \Theta, \Phi) \equiv f_{\text{DP}}(y | x; \Theta)$, a mixture over the bulk alone. A convenience interface to this bulk-only model is provided by the `dpmix()` wrapper.

```
R> fit <- dpmix(formula = y ~ x1 + x2 + x3, data = dat, backend = "crp",
+              kernel = "gamma", mcmc = mcmc_fixed)
```

Prior distributions for Θ and Φ can be specified by the user; defaults are weakly informative. For the DPM parameters Θ , the default prior employs stick-breaking priors for the mixture weights and conjugate priors for kernel parameters when conjugacy obtains. For the tail parameters Φ , the default adopts a regression-based specification for the threshold $u(x)$ and scale $\sigma(x)$, together with a weakly informative prior for the shape parameter ξ .

3.2. Statistical inference: estimation and prediction

Posterior simulation targets the joint posterior, which is given by

$$p(\Theta, \Phi | \mathcal{D}) \propto L(\Theta, \Phi; \mathcal{D}) p(\Theta, \Phi)$$

using the likelihood (5) and the prior structure described in Section 3.1. **CausalMixGPD** conducts MCMC via **nimble** (de Valpine *et al.* 2017, 2026). The fitted object supports `summary()` for posterior summaries, `plot()` for standard MCMC diagnostics, and `params()` for lightweight parameter extraction.

```
R> summary(fit) # posterior summaries for parameters and functionals
R> params(fit) # extract MCMC samples for parameters and functionals
R> plot(fit, family = "traceplot", params = "alpha")
R> # traceplot for concentration parameter
R> plot(fit, family = "auto") # automatic selection of diagnostic plots
```

The package provides posterior summaries for several functionals of the fitted conditional distribution defined in Section 2: the density $f(y | x)$, survival function $S(y | x) = 1 - F(y | x)$, quantile function $Q(\tau | x)$, and mean $E(y | x)$. For the `density`, `survival`, and `quantile` types, each functional is evaluated draw-by-draw from (2)–(4), and the retained draws are then summarized. For `type = "mean"`, the package employs analytical draw-level calculations whenever the kernel mean exists and, for spliced GPD models, the tail-shape parameter satisfies $\xi < 1$; if these conditions are not met, the ordinary mean is not reported and the user is directed to `type = "rmean"`. The restricted-mean target (`type = "rmean"`) provides a separate posterior predictive Monte Carlo calculation for a user-supplied cutoff and remains well-defined even when the ordinary mean does not exist. The following code illustrates how posterior density, quantile, and mean summaries are obtained at observed covariate values, together with credible intervals for quantiles and HPD intervals for means.

```
R> predict(fit, type = "density", interval = "credible", level = 0.95)
R> predict(fit, type = "quantile", index = c(0.25, 0.5, 0.75))
R> predict(fit, type = "mean", interval = "hpd", level = 0.9)
```

Predictions at new covariate values are obtained by calling `predict()` with a data set or covariate profiles sharing the same structure as the training data.

```
R> x_new <- as.data.frame(lapply(dat[, c("x1", "x2", "x3")], quantile,
+                               probs = c(0.25, 0.5, 0.75), na.rm = TRUE))
R> y_grid <- seq(0, 10, length.out = 200)
```

For a covariate profile $\mathbf{x}^*$, predictions are obtained from the posterior predictive density

$$p(y^* | \mathbf{x}^*, \mathcal{D}) = \int f(y^* | \mathbf{x}^*; \Theta, \Phi) p(\Theta, \Phi | \mathcal{D}) d\Theta d\Phi.$$

Predictions of the density and survival functions involve a grid of y values provided through the `y` argument, while predictions of quantiles require only the probability indices through the `index` argument. The code example that follows provides posterior survival and quantile predictions for new covariate profiles derived from the empirical quantiles of the predictors in the training data.

The argument-level controls for the above methods are provided in Appendix A. Specifically, Section A.1 extends the prior and backend definitions discussed previously, whereas Section A.2 summarizes the full prediction, interval, and diagnostic interfaces.

```
R> predict(fit, newdata = x_new, y = y_grid, type = "survival", level = 0.95)
R> predict(fit, newdata = x_new, type = "quantile", index = c(0.5, 0.99),
+          interval = "credible")
```

3.3. Clustering extension

```
R> data("nc_realX100_p3_k2", package = "CausalMixGPD")
R> dat_cl <- data.frame(y = nc_realX100_p3_k2$y, nc_realX100_p3_k2$X)
```

Under the clustering approach, the fitted mixture is regarded as a latent partition, and the output is cluster summaries. For a mixture model without the optional GPD tail, the function `dpmix.cluster()` is called to perform clustering; otherwise, `dpmgpd.cluster()` is used, with predictor dependence introduced via weights, mixture parameters, or both. Since labels assigned to components are arbitrary, the primary results include label-free summaries of the fitted mixture, namely the posterior similarity matrix (PSM), the representative partition, the number of clusters, and the cluster assignment scores (Binder 1978; Dahl 2006). The

same prediction interface `predict()` can produce cluster labels for training observations or assign clusters to new observations.

The code block below demonstrates how to fit a clustering model with covariates affecting both mixture weights and component parameters, by setting `type = "both"`.

```
R> fit_cluster <- dpmix.cluster(y ~ x1 + x2 + x3, data = dat_cl[1:90, ],
+                               kernel = "laplace", type = "both",
+                               components = 10, mcmc = mcmc_fixed)
```

The function `predict()` accepts two types of predictions: posterior similarity matrix or cluster labels, determined by the `type` argument.

```
R> predict(fit_cluster, type = "psm")
```

Cluster labels of new observations, along with corresponding assignment scores, are calculated as follows.

```
R> predict(fit_cluster, newdata = dat_cl[91:100, ], type = "label")
```

In Section 2.4, we mentioned that the software supports three types of dependence and evaluates clustering performance based on PSM summaries and representative partitions. Clustering-specific arguments are described in Appendix A.3.

3.4. Causal extension and propensity score augmentation

The causal inference workflow consists of estimating the causal framework separately per treatment arm and optionally augmenting the outcome covariates with a summary of the propensity score (PS). In the latter case (e.g., observational study), treatment allocation is assumed to be Bernoulli distributed, as follows,

$$A_i \mid \mathbf{x}_i \sim \text{Bernoulli}(e_\gamma(\mathbf{x}_i)),$$

and the predicted PS is included into the covariate vector to form an augmented regressor,

$$\mathbf{r}(\mathbf{x}) = (\mathbf{x}^\top, \psi(\hat{e}(\mathbf{x})))^\top,$$

where the PS link function (logit or probit) is determined by the PS parameter and the transformation $\psi(\cdot)$ used on the estimated propensity score depends on `ps_scale`; the default is the identity. Then, the outcome model in each arm is specified as a spliced DPM-GPD regression in the augmented covariates; the structure of splicing mixture regression is identical to that of Section 3.1. In the setting of randomized experiment, no PS estimation is needed, and the causal inference model for each arm uses the exact same modeling workflow as in the one-arm problem.

The code chunk below loads the data set generated from a benchmark setup and constructs the input `data.frame` required by the causal estimation workflows, consisting of the outcome, treatment variable, and covariates.

The invocation of `dpmgpd.causal()` trains spliced mixture regression for each treatment arm, with the causal model having the same structure as `dpmgpd()` but taking into account additional parameters controlling the PS model (`PS`, `ps_scale`, `ps_summary`) and different MCMC sampling strategies for the outcome and PS components (`mcmc_outcome`, `mcmc_ps`).

```
R> causal_fit <- dpmgpd.causal(formula = y ~ x1 + x2 + x3, data = causal_df,
+                              treat = "A", backend = "crp", kernel = "laplace",
+                              components = 5, PS = "logit", ps_scale = "logit",
+                              mcmc_outcome = mcmc_fixed, mcmc_ps = mcmc_fixed)
```


Estimated population-averaged quantities are computed using the functions `ate()` and `qte()`, and counterparts with treatment-population averaging are implemented through `att()` and `qtt()`. Confidence intervals are supported through interval estimation, and results can be visualized by invoking the generic `plot()` function with `type="effect"`, which plots estimated effects across quantiles along with the corresponding uncertainty bounds. The below code snippet shows the use of `qtt()` to compute treated population effects.

```
R> qtt_fit <- qtt(causal_fit, probs = c(0.5, 0.9, 0.95), interval = "credible")
R> summary(qtt_fit)
R> plot(qtt_fit, type = "effect")
```

Quantified conditional causal effects based on user-provided covariate profiles can be evaluated by the generic `predict()` function, which returns the selected summary of the posterior predictive outcome distribution depending on the value of the `type` parameter; the examples below produce survival probabilities at a grid of user-defined covariate values covering the empirical quartiles of the training variables.

```
R> predict(causal_fit, newdata = causal_xgrid, type = "survival",
+         y = rep(4, nrow(causal_xgrid)), id = seq_len(nrow(causal_xgrid)))
```

The corresponding bulk-only version is `dpmix.causal()`, where the causal estimand and prediction remain unchanged without the GPD tail. For randomized trials, the procedure is identical but without the PS adjustment, where `PS` can be set to `FALSE`. Other causal estimands, PS adjustments, and plots are documented in [Appendix A.3](#).

4. Data analysis I: clustering

This section demonstrates clustering under a bulk-only conditional mixture model. The goal is to discover latent subgroups among the observational units using the fitted mixture model, while accounting for partition uncertainty in a label-invariant manner. Unlike the density prediction focus of [Section 3](#), here the interest lies in partition summaries and cluster membership for out-of-sample data. Key outputs include the posterior similarity matrix (PSM), a representative partition, and cluster assignments for held-out observations.

4.1. Data and preprocessing

The Boston data set is used in this analysis, where each observation corresponds to a census tract. The response variable used is `medv`, representing median home prices in each area. Predictors used include the `lstat`, `rm`, and `nox` variables, which represent the level of neighborhood disadvantage, the size of the houses, and pollution levels, respectively. This combination of predictors constitutes a well-motivated choice for discovering representative clusters. We are interested in uncovering representative clusters from amongst the data, and predicting which latent clusters held-out census tracts belong to using out-of-sample prediction. An 80%/20% train-test split was conducted.

```
R> library("CausalMixGPD")
R> library("MASS")
R> set.seed(cmcpd_seed)
R> data("Boston", package = "MASS")
R> dat <- Boston
R> n <- nrow(dat)
R> idx_train <- sample(seq_len(n), size = floor(0.8 * n))
R> train_dat <- dat[idx_train, ]
R> test_dat <- dat[-idx_train, ]
```

Figure 4 summarizes the distribution of the Boston housing response, median home value (`medv`). The variability in the response is substantial, and the distribution exhibits a pronounced upper tail. These features motivate a flexible distributional model rather than relying only on a simple mean-regression specification.




Figure 4: Full-sample distribution of median house values (`medv`, in thousands of US dollars) for the 506 census tracts in the `MASS::Boston` data set. Left panel: boxplot with upper outliers highlighted. Right panel: density-scaled histogram on the same variable.

4.2. Model estimation and training cluster summaries

A bulk only conditional mixture model is applied to the `medv` dataset with three covariates. When `type = "weights"`, covariates determine the weights of the mixture while the parameters of the components are common across all covariates. This method is helpful for clustering where the prevalence of latent clusters may differ between neighborhoods.

```
R> .clust_rds <- file.path(manuscript_dir, "saved_fits", "fit_clust.rds")
R> if (file.exists(.clust_rds)) {
+   fit_clust <- readRDS(.clust_rds)
+ } else {
+   fit_clust <- dpmix.cluster(formula = medv ~ ., data = train_dat,
+                             kernel = "normal", type = "weights",
+                             components = 10, mcmc = mcmc_fixed)
+   dir.create(dirname(.clust_rds), showWarnings = FALSE, recursive = TRUE)
+   saveRDS(fit_clust, .clust_rds)
+ }
```

Since mixture components cannot be identified up to permutation, clusters in the training data are represented using the posterior similarity matrix and sample training labels. The posterior similarity matrix gives the main label-invariant uncertainty summary by recording posterior clustering probabilities.

```
R> z_train_psm <- predict(fit_clust, type = "psm")
```

```
R> plot(z_train_psm, type = "summary")
```

As seen from Figure 5, the PSM reveals clear block structure along the diagonal, implying that the model can detect multiple sets of tracts whose clustering probability is relatively high. While there are clear separations, there are still some uncertainties for those tracts on the edge of two or more neighboring groups, which reflects reasonable posterior uncertainty in a non-experimental housing dataset, since features of neighborhoods tend to change gradually rather than abruptly.

```
R> z_train_lab <- predict(fit_clust, type = "label")
```




Figure 5: Posterior similarity matrix (PSM) summary for the Boston training sample, returned by `plot(predict(fit_clust, type = "psm"), type = "summary")`. Each cell (i, ℓ) encodes the posterior clustering probability $\Pr(z_i = z_\ell \mid \text{data})$, with rows and columns ordered to expose block structure. Darker shading corresponds to higher clustering probability.

Figure 6 shows the representative training cluster summary produced by the package `plot()` method applied to the training labels object.

```
R> plot(z_train_lab, type = "summary")
```




Figure 6: Representative training-cluster summary for the Boston `medv` analysis, returned by the `plot()` method applied to the labels object from `predict(fit_clust, type = "label")`. Training observations are shown on the response scale and coloured by their assigned cluster in the representative partition.

The representative partition translates the posterior clustering structure into an interpretable summary on the response scale. The inferred groups separate tracts primarily by housing-value level, but they also reflect broader differences in neighborhood composition and environmental quality. Thus, the clustering output is not merely a numerical partition induced by the mixture fit; it yields a meaningful summary of heterogeneity in the Boston housing data.

4.3. Prediction for new observations

Cluster assignment for new observations is performed through the same prediction interface.

The output is a label in the space of the representative training clusters, obtained by aggregating posterior predictive clustering evidence at the cluster level.

```
R> z_test <- predict(fit_clust, newdata = test_dat, type = "label")
```

Figure 7 displays the distribution of predicted cluster assignments across the test sample, produced by the package `plot()` method applied to the test labels object. The printed summary below gives the median response and covariate profiles for each predicted cluster in the test set, returned by `cluster_profiles(predict(fit_clust, newdata = test_dat, type = "label"))`.

```
R> cluster_profiles(z_test)
```

	cluster	n	medv_mean	medv_sd	crim_mean	crim_sd	zn_mean			
1	C2	46	20.87174	3.133309	0.7925048	1.725793	2.902174			
2	C3	30	15.36000	7.796666	13.5420153	18.960580	2.666667			
3	C1	26	31.68846	9.695497	1.0496638	1.752975	22.615385			
			zn_sd	indus_mean	indus_sd	chas_mean	chas_sd	nox_mean		
1			7.856787	11.349783	6.077225	0.06521739	0.2496374	0.5286957		
2			14.605935	18.046333	2.961173	0.03333333	0.1825742	0.6751667		
3			27.368708	7.341538	6.405760	0.07692308	0.2717465	0.5259231		
				nox_sd	rm_mean	rm_sd	age_mean	age_sd	dis_mean	
1				0.08399335	6.061370	0.3964721	67.79565	28.14555	3.868787	
2				0.07969212	6.010900	0.6312801	89.93000	15.24594	2.206400	
3				0.11591063	6.884615	1.0538402	58.46923	28.94230	4.199923	
					dis_sd	rad_mean	rad_sd	tax_mean	tax_sd	ptratio_mean
1					1.630337	5.065217	4.291605	319.1087	97.58921	18.43261
2					1.418189	20.600000	7.753086	622.6333	100.56272	19.99000
3					2.196438	8.076923	8.153149	382.9615	150.14885	16.91154
						ptratio_sd	black_mean	black_sd	lstat_mean	lstat_sd
1						1.937244	380.0700	34.36248	11.606522	3.770490
2						1.264461	277.3020	156.42949	19.653000	6.326669
3						2.701307	387.8165	12.37760	6.598462	2.676518

```
R> plot(z_test, type = "sizes")
```




Figure 7: Bar plot of predicted cluster assignments for the held-out Boston test sample (the 20% of `medv` observations excluded from training), returned by `plot(predict(fit_clust, newdata = test_dat, type = "label"), type = "sizes")`. Each bar corresponds to one representative training cluster from Figure 6; bar heights give the number of test tracts assigned to that cluster.

The obtained profiles for the clusters correspond to our expectations concerning the representative partition. Cluster C1 includes tracts with higher average house values (`medv`), more

rooms (**rm**), greater residential zoning (**zn**), and lower pollution (**nox**). On the contrary, cluster C3 represents the poorer area with fewer rooms (**rm**), low house value (**medv**), high rate of crime (**crim**) and industrialization (**indus**), and the highest level of average pollution (**nox**). Cluster C2 represents an area in between. For test data, clusters characterized by larger average house prices (**medv**) have fewer neighborhood disadvantages such as lower crime per capita (**crim**), more spacious houses (**rm**), and less pollution (**nox**); poorer clusters represent the opposite pattern.

The presented case study demonstrates the usage of a representative partition and its cluster description as well as how we assign posterior predictive membership probabilities for a new observation. The use of representative clusters makes the connection between partition uncertainty and distribution of the response given covariates clearer.

5. Data analysis II: causal inference

This section describes the estimands from the two-arm spliced density regression. In contrast to the partition estimation of Section 4, the focus here is on mean and quantile differences for both unconditional and conditional causal estimands.

5.1. Data and estimands

To estimate unconditional and conditional causal estimands using the proposed approach, the **lalonge** dataset from the **MatchIt** R package is used for benchmarking. The outcome is the 1978 annual income level measured by **re78** (\$), and the treatment indicator corresponds to program participation (variable **treat**). Covariates include age, years of schooling, race, marriage status, college degree status, and prior earnings **re74** and **re75**. Earnings typically exhibit considerable skewness, concentrate near zero, and have a pronounced upper tail, motivating the use of quantile and tail summaries rather than the mean alone.

The six estimands in Table 1 are computed following Section 2.5. Standardized estimands (ATE, QTE) are averaged across the empirical distribution of $\{\mathbf{x}_i\}_{i=1}^n$, while conditional estimands (CATE, CQTE) are evaluated at selected covariate profiles.

The data are first loaded and the design matrix constructed.

```
R> data("lalonge", package = "MatchIt")
R> app_lalonge_covars <- within(lalonge, {
+   race <- factor(race, levels = c("white", "black", "hispan"))
+   married <- factor(married, levels = c(0, 1), labels = c("no", "yes"))
+   nodegree <- factor(nodegree, levels = c(0, 1), labels = c("no", "yes"))
+ })
R> app_lalonge_scale_vars <- c("age", "educ", "re74", "re75")
R> app_lalonge_scale_center <- vapply(
+   app_lalonge_scale_vars,
+   function(v) mean(app_lalonge_covars[[v]], na.rm = TRUE),
+   numeric(1)
+ )
R> app_lalonge_scale_scale <- vapply(
+   app_lalonge_scale_vars,
+   function(v) stats::sd(app_lalonge_covars[[v]], na.rm = TRUE),
+   numeric(1)
+ )
R> for (v in app_lalonge_scale_vars) {
+   app_lalonge_covars[[paste0(v, "_z")]] <-
+     (app_lalonge_covars[[v]] - app_lalonge_scale_center[[v]]) /
+     app_lalonge_scale_scale[[v]]
+ }
R> app_lalonge_x_formula <- ~age_z + educ_z + race + married + nodegree + re74_z +
+   re75_z
R> app_lalonge_y <- (app_lalonge_covars$re78 + 0.5)/1000
```

```
R> app_lalonde_A <- as.integer(app_lalonde_covars$treat)
R> app_lalonde_taus <- c(0.25, 0.5, 0.75, 0.9, 0.95)
R> app_lalonde_X <- stats::model.matrix(
+   app_lalonde_x_formula,
+   data = app_lalonde_covars)[, -1, drop = FALSE]
```

Figure 8 provides full-sample descriptive views of observed 1978 earnings (`re78`) by treatment assignment. Both groups exhibit substantial mass near zero together with pronounced right-skewness, and the arm-specific distributions differ in ways that are not well summarized by a simple location shift. These features motivate modeling the full treatment-specific outcome distributions directly rather than relying solely on average effects.




Figure 8: Full-sample distribution of 1978 annual earnings (`re78`, in US dollars) from the `MatchIt::lalonde` data set, stratified by treatment arm. Left panel: arm-specific boxplots with upper outliers highlighted. Right panel: overlaid density-scaled histograms for the control and treated arms.

5.2. Model and fit

Arm-specific conditional outcome distributions $F_a(\cdot | \mathbf{x})$ are modeled via the spliced bulk–tail construction described in Section 2.3:

$$f_a(y | \mathbf{x}) = \sum_j w_{aj}(\mathbf{x}) k(y | \Theta_{aj}) \mathbf{1}(y \leq u) + \{1 - F_a(u | \mathbf{x})\} g(y | u, \sigma_a, \xi_a) \mathbf{1}(y > u).$$

The bulk kernel $k(\cdot)$ represents a Dirichlet process mixture regression with a Chinese restaurant process backbone, while the tail density $g(\cdot)$ follows the GPD distribution with positive scale $\sigma_a > 0$ and tail index ξ_a . In our model, the GPD tail index ξ_a regulates tail heaviness (e.g., if $\xi_a > 0$, it indicates a heavier tail). Our model mitigates poor finite-sample precision for extreme quantiles through nonparametric bulk and EVT-consistent tail modeling.

Here, a two-sample mixture is fitted with the gamma bulk kernel and the same MCMC configuration throughout. The aim is not to draw definitive conclusions but to demonstrate how the package produces standardized and conditional earnings contrasts under the spliced DPM-GPD framework.

```
R> fit <- dpmgpd.causal(y = app_lalonde_y, X = app_lalonde_X,
+   treat = app_lalonde_A, backend = "crp",
+   kernel = "gamma", components = 10, mcmc = mcmc_fixed)
```

5.3. Posterior summaries

The fitted causal object defines two arm-specific outcome distributions from which all reported summaries are derived. The framework is therefore not limited to a single regression

coefficient or scalar contrast; mean, quantile, and tail-sensitive causal estimands are all coherent functionals of the same fitted posterior distribution. The printed output below reports wall-clock MCMC run times for each arm when timing information is available.

```
CausalMixGPD causal fit
PS model: Bayesian logit (A | X)
Outcome (treated): backend = crp | kernel = gamma
Outcome (control): backend = crp | kernel = gamma
GPD tail (treated/control): TRUE / TRUE
Timing (sec): total = NA | PS = NA | control = NA | treated = NA
```

5.4. Standardized estimands: ATE and QTE

The overall treatment effect is summarized with ATE and QTE.

```
R> ate_fit <- ate(fit, interval = "hpd", level = 0.95)
R> summary(ate_fit)
```

```
ATE Summary
=====
Prediction points: 1
Conditional: NO | PS used: YES
Interval: hpd

Model specification:
  Backend (trt/con): crp / crp
  Kernel (trt/con): gamma / gamma
  GPD tail (trt/con): YES / YES

ATE estimates:
  estimate lower upper
    -2.543 -5.341 0.396
```

The posterior ATE estimate is negative on the scaled earnings scale, but the corresponding interval includes zero. Thus, after averaging over the empirical covariate distribution, the analysis does not provide strong evidence of a uniform overall mean effect of the training program.

```
R> qte_fit <- qte(fit, probs = c(0.25, 0.5, 0.75), interval = "credible")
```

Figure 9 displays two standardized QTE views side by side from the causal model: the estimated QTE profile across the evaluated quantiles and the treated-versus-control quantile summaries with uncertainty intervals.

```
R> qte_effect_plot <- plot(qte_fit, type = "effect")
R> qte_arms_plot <- plot(qte_fit, type = "arms")
R> qte_effect_plot + qte_arms_plot + plot_layout(ncol = 2)
```

Estimates of QTEs are all negative for the quantiles analyzed. However, the confidence interval for the median QTE and the for the 3rd quartile do not contain zero, indicating that the training program have a negative impact on high earners.

5.5. Conditional estimands for representative profiles

To characterize heterogeneity in the treatment effect, conditional estimands are computed at four representative covariate profiles spanning the observed covariate range. Profile values are expressed on the original data scale, using empirical quantiles for continuous variables and observed levels for categorical variables. Continuous predictors are standardized using training-set means and standard deviations before being passed to the model.




Figure 9: Standardized quantile treatment effect (QTE) summaries for Lalonde 1978 earnings, returned by `qte(fit, probs = c(0.25, 0.50, 0.75), interval = "credible")`. Left panel: effect curve $Q_1^m(\tau) - Q_0^m(\tau)$ versus τ with pointwise 95% credible band. Right panel: standardized treated and control quantile curves $Q_1^m(\tau)$ and $Q_0^m(\tau)$ on the original earnings scale, with pointwise credible bands.

Table 4: Representative Lalonde participant profiles used as new-data inputs for the conditional causal contrasts reported in Figure 10 and in the printed CQTE summary below. Columns (Profile 1–Profile 4) give the covariate values for one synthetic individual on the original data scale: **age** (years), **educ** (years of schooling), **race**, **married**, **nodegree**, and prior annual earnings **re74** and **re75** (US dollars).

	Profile 1	Profile 2	Profile 3	Profile 4
age	20.0000	25.0000	32.0000	43.0000
educ	9.0000	11.0000	12.0000	13.0000
race	black	hispan	white	black
married	no	no	yes	yes
nodegree	yes	no	no	yes
re74	0.0000	47.4143	2705.7476	9383.2356
re75	0.0000	16.4710	1389.6486	4201.8870

```
R> cate_fit <- cate(fit, newdata = xnew, interval = "hpd")
R> summary(cate_fit)
```

CATE Summary

```
=====
Prediction points: 4
Conditional: YES | PS used: YES
Interval: hpd
```

Model specification:

```
Backend (trt/con): crp / crp
Kernel (trt/con): gamma / gamma
GPD tail (trt/con): YES / YES
```

CATE estimates:

```
profile estimate lower upper
Profile 1 -1.978 -5.481 1.448
Profile 2 -2.685 -7.837 1.612
Profile 3 -1.744 -7.203 3.256
Profile 4 -7.725 -31.332 15.31
```

Figure 10 shows the CATE summaries across the four representative profiles. None of the four profiles show a statistically significant CATE, as all 95% HPD intervals




Figure 10: Profile-specific conditional average treatment effects for Lalonde 1978 earnings at the four representative covariate profiles (Profile 1–Profile 4) of Table 4, returned by `cate(fit, newdata = xnew, interval = "hpd")`. Points denote posterior estimates of $\text{CATE}(x) = \mu_1(x) - \mu_0(x)$ and vertical bars give 95% HPD intervals.

include zero. However, the point estimates suggest that the program may have a negative effect on earnings for Profiles 4, while the effects for Profiles 1, 2 and 4 are closer to zero.

```
R> cqte_fit <- cqte(fit, probs = c(0.25, 0.5, 0.75, 0.9, 0.95),
+                  newdata = xnew, interval = "credible")
R> summary(cqte_fit)$effect_table # full CQTE summary table
```

	profile	tau	estimate	lower	upper
1	Profile 1	0.25	-2.521841198	-7.744296	4.0271372
2	Profile 1	0.50	-3.717253897	-7.385001	-0.6244419
3	Profile 1	0.75	-3.036742586	-7.347654	1.4575352
4	Profile 1	0.90	-1.616077528	-7.876948	5.9154540
5	Profile 1	0.95	-0.002597363	-8.645220	10.5596639
6	Profile 2	0.25	-1.388753348	-8.374172	3.5607329
7	Profile 2	0.50	-4.461346476	-9.850380	-0.1311481
8	Profile 2	0.75	-4.406719223	-10.769044	0.5638009
9	Profile 2	0.90	-3.966523104	-11.325236	3.6031938
10	Profile 2	0.95	-3.546162922	-13.181625	6.9877301
11	Profile 3	0.25	-1.400116612	-6.699770	2.7939363
12	Profile 3	0.50	-1.800035981	-7.721600	3.5216058
13	Profile 3	0.75	-2.604699710	-8.484734	4.2098620
14	Profile 3	0.90	-2.871274028	-8.828207	4.2015030
15	Profile 3	0.95	-2.528672343	-10.414719	8.3735820
16	Profile 4	0.25	-2.553434640	-9.052905	4.3991099
17	Profile 4	0.50	-5.901260763	-17.654505	6.3991686
18	Profile 4	0.75	-11.486181749	-45.838709	20.1632763
19	Profile 4	0.90	-19.094601788	-86.637680	42.1730175
20	Profile 4	0.95	-24.260813769	-119.477696	64.8149038

The CQTE estimates reveal heterogeneity across all covariate profiles and quantile levels. At low and median quantile levels, the estimated program effect is negative or close to zero, with the median CQTE significantly below zero for Profiles 1 and 2. Uncertainty increases sharply at the upper quantiles, especially for Profile 4, because upper-quantile contrasts rely on fewer observations and are more sensitive to the spliced tail. Overall, the CQTE table provides no clear evidence of positive program effects: upper-tail intervals sometimes extend into positive values, but the corresponding point estimates remain negative.

5.6. Results and interpretation

Two questions bear on the interpretation: the overall program effect on earnings, and heterogeneity of effects across participant profiles. The first yields no statistically compelling evidence: the average treatment effect is close to zero and the estimated QTE are negative at lower and mid-quantiles. The second indicates heterogeneity in magnitude and uncertainty across profiles, but not a clear reversal in sign: CQTE point estimates remain nonpositive across the reported quantiles, while upper-tail intervals widen substantially for some profiles. This analysis underscores the value of quantile-based methods for heavy-tailed outcomes. A mean-only analysis provides little evidence of a program effect; the QTE, by contrast, reveals heterogeneity in the treatment effect across the outcome distribution.

6. Discussion

CausalMixGPD combines flexible bulk density estimation, optional EVT-based tail extrapolation, predictor-dependent clustering, and quantile- and density-level causal inference within one package. The same fitted model supports posterior summaries for densities, survival summaries, analytical means (when they exist), restricted means, quantiles, and treatment contrasts. This is useful in settings where heavy tails and density-level effects are both of concern.

The package is considerably broader than the examples emphasized in this article. Users may work with bulk-only or spliced models, multiple kernel families on the real line or on positive support, stick-breaking or CRP backends, and clustering or causal wrappers built upon the same predictive machinery. This shared infrastructure lets users move from exploratory subgroup analysis to tail-focused prediction and causal reporting without changing the underlying modeling language.

Several practical limitations warrant acknowledgment. First, computation can be demanding. Mixture models with covariate dependence, optional tail components, and arm-specific causal fits are inherently MCMC-intensive, and costs escalate further with sensitivity analyses, longer chains, or richer predictor sets.

Second, model tuning requires attention. A spliced analysis inherits specification choices from both the DPM and peaks-over-threshold components, including kernel family, truncation level, threshold specification, and the number of exceedances available for tail learning. When overlap is weak or exceedances are sparse, posterior uncertainty can be misleading; for clustering, the posterior similarity matrix is often more informative than any single partition summary.

Third, the current scope is circumscribed. The causal implementation is restricted to binary treatments and continuous outcomes under standard consistency, ignorability, and overlap assumptions. Propensity-score augmentation can improve covariate adjustment in observational studies, but it is not a substitute for careful study design, overlap diagnostics, or sensitivity analysis for unmeasured confounding.

Finally, this article treats **CausalMixGPD** as a software contribution and describes the implemented modeling workflow self-containedly. Formal methodological development beyond the software and examples presented here should be evaluated separately from the package interface, documentation, and replication material.

7. Software Availability

Version 0.8.0 of the **CausalMixGPD** R package is the version used for this article. The package is licensed under GPL-3 and is available from CRAN. The development version is available from the GitHub repository <https://github.com/arnabaich96/CausalMixGPD>.

References

- Aich A (2026a). *CausalMixGPD: Bayesian Nonparametric Conditional Density Modeling in Causal Inference and Clustering with a Heavy-Tail Extension*. doi:10.32614/CRAN.package.CausalMixGPD. R package version 0.8.0, URL <https://cran.r-project.org/package=CausalMixGPD>.
- Aich A (2026b). “A Semiparametric Bayesian Approach for Estimating Extreme Conditional Quantile Treatment Effect for Heavy-tailed Data.” doi:10.5281/zenodo.19672760. Zenodo record, URL <https://zenodo.org/records/19672760>.
- Balkema AA, de Haan L (1974). “Residual Life Time at Great Age.” *The Annals of Probability*, **2**(5), 792–804. doi:10.1214/aop/1176996548.
- Behrens CN, Lopes HF, Gamerman D (2004). “Bayesian Analysis of Extreme Events with Threshold Estimation.” *Statistical Modelling*, **4**(3), 227–244. doi:10.1191/1471082X04st075oa.
- Beraha M, Guindani B, Gianella M, Guglielmi A (2025). “BayesMix: Bayesian Mixture Models in C++.” *Journal of Statistical Software*, **112**(9), 1–41. doi:10.18637/jss.v112.i09.
- Binder DA (1978). “Bayesian Cluster Analysis.” *Biometrika*, **65**(1), 31–38. doi:10.1093/biomet/65.1.31.
- Burns N, Daniels MJ (2023). “Truncation Approximation for Enriched Dirichlet Process Mixture Models.” doi:10.48550/arXiv.2305.01631. 2305.01631, URL <https://arxiv.org/abs/2305.01631>.
- Coles S (2001). *An Introduction to Statistical Modeling of Extreme Values*. Springer, London. doi:10.1007/978-1-4471-3675-0.
- Corradin R, Canale A, Nipoti B (2021). “BNPmix: An R Package for Bayesian Nonparametric Modeling via Pitman–Yor Mixtures.” *Journal of Statistical Software*, **100**(15), 1–47. doi:10.18637/jss.v100.i15.
- Dahl DB (2006). “Model-Based Clustering for Expression Data via a Dirichlet Process Mixture Model.” In KA Do, P Müller, M Vannucci (eds.), *Bayesian Inference for Gene Expression and Proteomics*, pp. 201–218. Cambridge University Press, Cambridge. doi:10.1017/CB09780511584589.011. URL <https://doi.org/10.1017/CB09780511584589.011>.
- De Iorio M, Müller P, Rosner GL, MacEachern SN (2004). “An ANOVA Model for Dependent Random Measures.” *Journal of the American Statistical Association*, **99**(465), 205–215. doi:10.1198/016214504000000205.
- de Valpine P, Paciorek C, Turek D, Michaud N, Anderson-Bergman C, Obermeyer F, Wehrhahn Cortes C, Rodríguez A, Temple Lang D, Zhang W, Paganin S, Hug J, van Dam-Bates P (2026). *nimble: MCMC, Particle Filtering, and Programmable Hierarchical Modeling*. doi:10.32614/CRAN.package.nimble. R package on CRAN, URL <https://cran.r-project.org/package=nimble>.
- de Valpine P, Turek D, Paciorek C, Anderson-Bergman C, Temple Lang D, Bodik R (2017). “Programming with Models: Writing Statistical Algorithms for General Model Structures with NIMBLE.” *Journal of Computational and Graphical Statistics*, **26**(2), 403–413. doi:10.1080/10618600.2016.1172487.
- Dinari O, Yu A, Freifeld O, Fisher JW (2019). “Distributed MCMC Inference in Dirichlet Process Mixture Models Using Julia.” In *2019 19th IEEE/ACM International Symposium*

- on Cluster, Cloud and Grid Computing (CCGRID), pp. 518–525. doi:10.1109/CCGRID.2019.00066.
- Dorie V, Hill J (2025). *bartCause: Causal Inference Using Bayesian Additive Regression Trees*. doi:10.32614/CRAN.package.bartCause. R package on CRAN, URL <https://cran.r-project.org/package=bartCause>.
- DoWhy Contributors (2026). “DoWhy: An End-to-End Library for Causal Inference.” Python software; accessed 2026-05-07, URL <https://github.com/py-why/dowhy>.
- EconML Contributors (2026). “EconML: A Python Package for Machine Learning-Based Heterogeneous Treatment Effects.” Python software; accessed 2026-05-07, URL <https://github.com/py-why/EconML>.
- Escobar MD, West M (1995). “Bayesian Density Estimation and Inference Using Mixtures.” *Journal of the American Statistical Association*, **90**(430), 577–588. doi:10.1080/01621459.1995.10476550.
- Ferguson TS (1973). “A Bayesian Analysis of Some Nonparametric Problems.” *The Annals of Statistics*, **1**(2), 209–230. doi:10.1214/aos/1176342360.
- Frigessi A, Haug O, Rue H (2002). “A Dynamic Mixture Model for Unsupervised Tail Estimation without Threshold Selection.” *Extremes*, **5**(3), 219–235. doi:10.1023/A:1024072610684.
- Hahn PR, Murray JS, Carvalho CM (2020). “Bayesian Regression Tree Models for Causal Inference: Regularization, Confounding, and Heterogeneous Effects.” *Bayesian Analysis*, **15**(3), 965–1056. doi:10.1214/19-BA1195. URL <https://projecteuclid.org/journals/bayesian-analysis/volume-15/issue-3/Bayesian-Regression-Tree-Models-for-Causal-Inference--Regularization-Confounding/10.1214/19-BA1195.full>.
- Hill JL (2011). “Bayesian Nonparametric Modeling for Causal Inference.” *Journal of Computational and Graphical Statistics*, **20**(1), 217–240. doi:10.1198/jcgs.2010.08162.
- Hu Y, Scarrott C (2018). “evmix: An R Package for Extreme Value Mixture Modeling, Threshold Estimation and Boundary Corrected Kernel Density Estimation.” *Journal of Statistical Software*, **84**(5), 1–27. doi:10.18637/jss.v084.i05.
- Imbens GW, Rubin DB (2015). *Causal Inference for Statistics, Social, and Biomedical Sciences: An Introduction*. Cambridge University Press, Cambridge. doi:10.1017/CB09781139025751.
- Liverani S, Hastie DI, Azizi L, Papathomas M, Richardson S (2015). “PReMiuM: An R Package for Profile Regression Mixture Models Using Dirichlet Processes.” *Journal of Statistical Software*, **64**(7), 1–30. doi:10.18637/jss.v064.i07.
- MacEachern SN (1999). *Dependent Nonparametric Processes*. ASA Proceedings of the Section on Bayesian Statistical Science, URL <https://u.osu.edu/maceachern.1/files/2025/05/1999-MacEachern-JSM-Proceedings.pdf>.
- MacEachern SN (2000). *Dependent Dirichlet Processes*. Technical Report, Department of Statistics, The Ohio State University, URL <https://u.osu.edu/maceachern.1/files/2025/05/2000-MacEachern-DDP-Tech-Report.pdf>.
- Murray JS, Hahn PR, Carvalho C (2024). *bcf: Causal Inference Using Bayesian Causal Forests*. doi:10.32614/CRAN.package.bcf. R package on CRAN, URL <https://cran.r-project.org/package=bcf>.

- Neal RM (2000). “Markov Chain Sampling Methods for Dirichlet Process Mixture Models.” *Journal of Computational and Graphical Statistics*, **9**(2), 249–265. doi:[10.1080/10618600.2000.10474879](https://doi.org/10.1080/10618600.2000.10474879).
- Pedregosa F, Varoquaux G, Gramfort A, Michel V, Thirion B, Grisel O, Blondel M, Prettenhofer P, Weiss R, Dubourg V, Vanderplas J, Passos A, Cournapeau D, Brucher M, Perrot M, Duchesnay É (2011). “Scikit-Learn: Machine Learning in Python.” *Journal of Machine Learning Research*, **12**, 2825–2830. URL <https://jmlr.org/papers/v12/pedregosa11a.html>.
- Pickands J (1975). “Statistical Inference Using Extreme Order Statistics.” *The Annals of Statistics*, **3**(1), 119–131. doi:[10.1214/aos/1176343003](https://doi.org/10.1214/aos/1176343003).
- Quintana FA, Müller P, Jara A, MacEachern SN (2022). “The Dependent Dirichlet Process and Related Models.” *Statistical Science*, **37**(1), 24–49. doi:[10.1214/20-STS819](https://doi.org/10.1214/20-STS819).
- Ren L, Du L, Carin L, Dunson DB (2011). “The Logistic Stick-Breaking Process.” *Journal of Machine Learning Research*, **12**(7), 203–239. URL <https://www.jmlr.org/papers/v12/ren11a.html>.
- Ribatet M, Dutang C (2024). *POT: Generalized Pareto Distribution and Peaks Over Threshold*. doi:[10.32614/CRAN.package.POT](https://doi.org/10.32614/CRAN.package.POT). R package on CRAN, URL <https://cran.r-project.org/package=POT>.
- Rosenbaum PR, Rubin DB (1983). “The Central Role of the Propensity Score in Observational Studies for Causal Effects.” *Biometrika*, **70**(1), 41–55. doi:[10.1093/biomet/70.1.41](https://doi.org/10.1093/biomet/70.1.41).
- Ross GJ, Markwick D (2023). *dirichletprocess: Build Dirichlet Process Objects for Bayesian Modelling*. doi:[10.32614/CRAN.package.dirichletprocess](https://doi.org/10.32614/CRAN.package.dirichletprocess). R package on CRAN, URL <https://cran.r-project.org/package=dirichletprocess>.
- Rubin DB (1974). “Estimating Causal Effects of Treatments in Randomized and Nonrandomized Studies.” *Journal of Educational Psychology*, **66**(5), 688–701. doi:[10.1037/h0037350](https://doi.org/10.1037/h0037350).
- Selva F, Fuentes-García R, Gil-Leyva MF (2025). “pyrichlet: A Python Package for Density Estimation and Clustering Using Gaussian Mixture Models.” *Journal of Statistical Software*, **112**(8), 1–39. doi:[10.18637/jss.v112.i08](https://doi.org/10.18637/jss.v112.i08).
- Sethuraman J (1994). “A Constructive Definition of Dirichlet Priors.” *Statistica Sinica*, **4**(2), 639–650. URL <https://www3.stat.sinica.edu.tw/statistica/j4n2/j4n216/j4n216.htm>.
- Wade S, Dunson DB, Petrone S, Trippa L (2014). “Improving Prediction from Dirichlet Process Mixtures via Enrichment.” *Journal of Machine Learning Research*, **15**(30), 1041–1071. URL <https://jmlr.org/papers/v15/wade14a.html>.

A. Customization of package and additional features

This appendix details the customizable features of **CausalMixGPD**. It covers argument-level customization for model fitting, posterior manipulation, and workflow-specific extensions.

A.1. Model fitting options and prior specifications

Besides the customizable features illustrated in Sections 2.1, 2.2, 2.5, and 3.1, this section lists arguments that allow kernel choice, latent structure, and prior customization alternative to default choices.

- **Kernel families and support:** the argument `kernel` selects one of seven kernel families on either the positive line or the real line. Positive-support options include "gamma" and "lognormal". They also include "invgauss" and "amoroso". Real-line options are "normal", "laplace", and "cauchy". For causal models, `kernel` can be a single string or a length-two vector specifying different DPM kernels for the two arms. The choice of kernel should match the support of the outcome variable.
- **Backend type and truncation rule:** the argument `backend` determines the structure of the mixing coefficients. The value `sb` stands for the stick-breaking scheme while `crp` denotes the Chinese Restaurant Process latent allocation scheme. For causal models, `backend` can be either a single string or a length-two vector specifying different latent allocation schemes for the two arms. The argument `components` sets the maximum number of mixture components (defaulting to the sample size). The argument `epsilon` specifies a truncation threshold below which components are dropped and the smallest weight is readjusted. Both arguments can be customized separately for each causal arm.
- **Parameter specification and prior setting:** the argument `param_specs` is a named list giving the parameters of the kernel or tail distribution. Each kernel/tail parameter supports three modes: "fixed", "dist", and "link". With "dist", the prior can be chosen as any distribution compatible with the parameter and supported by **nimble**, with hyperparameters specified in the default registry. With "link", the kernel/tail parameter is related to the covariates through a linear regression using one of five link functions ("identity", "exp", "log", "softplus", or "power"); the link power is adjusted via `link_power` when the last option is selected. With "fixed", the parameter is fixed to the constant value specified in `param_specs`. For causal models, `param_specs` can be either a single named list or a length-two list specifying different parameter specifications for the two arms.

A.2. Fitting, prediction, and posterior output

This section gathers the arguments used after a model class has been chosen: sampler configuration, monitoring, and fitting interfaces.

- **MCMC configuration:** `mcmc` controls `niter`, `nburnin`, `thin`, `nchains`, and `seed`; causal models split this into `mcmc_outcome` and `mcmc_ps`. Runtime controls include `show_progress`, `quiet`, `parallel_chains` with `workers`, `parallel_arms` for arm-wise causal fitting, and `timing`. CRP-based backends additionally use `z_update_every` to tune the frequency of global label updates.
- **Monitoring and fitting interfaces:** `monitor = "core"` retains the main structural parameters, whereas `"full"` retains all nodes. `monitor_latent = TRUE` stores allocation labels $z_1, \dots, z_n$, and `monitor_v = TRUE` stores stick-breaking fractions for SB models. Models can be supplied through a formula interface (`formula`, `data`) or directly through `y` and `X`; causal fits additionally use `treat`.

A.3. Workflow-specific extensions

This section records controls specific to the clustering and causal workflows in Sections 3.3, 3.4, 4, and 5. Representative outputs appear in Figures 6, 7, 9, and 10.

- **Propensity score customization:** observational causal fits can estimate propensity scores with `PS = "logit", "probit", "naive", or FALSE`. The estimated score is appended to the outcome model on either the logit scale, `ps_scale = "logit"`, or the probability scale, `"prob"`. Posterior summaries use `ps_summary = "mean" or "median"`. Numerical stability is controlled through `ps_clamp`, shrinkage through `ps_prior`, and the two-stage fitting scheme through separate `mcmc_ps` and `mcmc_outcome` lists.
- **Clustering dependence mode:** clustering fits accept `type = "weights", "param", or "both"` to place predictor dependence in the mixture weights, kernel parameters, or both. The `newdata` argument controls whether label assignment is restricted to training observations or extended to additional covariate profiles.

A.4. Utilities and computing support

The last appendix section includes computing and data resources that accompany the main modeling process.

- **Simulators and bundled datasets:** the package includes reproducible benchmark generators for bulk, tail, and causal analyses. Random-number seeds are supplied by the `seed` argument. The package also bundles twelve datasets spanning unconditional, conditional, and causal settings on positive and real-line supports, with and without tail components.
- **Computational considerations:** large-scale analyses can use `parallel_chains` with `workers`, `parallel_arms` for causal fits, `chunk_size` for prediction in parallel chunks, and `ndraws_pred` for sampling from posterior predictive parameter distributions. The `show_progress` and `quiet` arguments control progress reporting and compiler verbosity, while `timing = TRUE` records runtime summaries.

Affiliation:

Arnab Aich
 Department of Statistics
 Florida State University
 Tallahassee, FL, USA
 E-mail: aaich@fsu.edu
 ORCID: [0009-0005-7801-6701](https://orcid.org/0009-0005-7801-6701)

Indrabati Bhattacharya
 Department of Statistics
 Florida State University
 Tallahassee, FL, USA
 E-mail: ib22g@fsu.edu
 ORCID: [0000-0003-4567-0027](https://orcid.org/0000-0003-4567-0027)