Journal of Statistical Software

MMMMMM YYYY, Volume VV, Issue II. doi: 10.18637/jss.v000.i00

DPmixGPD: Bayesian Conditional Dirichlet Process Mixtures with Generalized Pareto Tails and Causal Extensions

Arnab Aich
Florida State University

Abstract

Heavy-tailed outcomes arise in many scientific domains where rare events dominate risk, cost, or policy decisions. Classical parametric models often fit the center of the distribution at the expense of tail misspecification, while purely nonparametric methods typically lack principled extrapolation in regions with sparse data. Extreme value theory (EVT) provides asymptotic justification for generalized Pareto distributions (GPDs) for threshold exceedances, whereas Bayesian nonparametric Dirichlet process mixtures (DPMs) provide flexible modeling for heterogeneous bulk behavior. This article introduces `DPmixGPD`, an R package implementing a Bayesian semiparametric framework that combines conditional DPMs for the bulk distribution with covariate-dependent GPD tails for exceedances, with posterior predictive inference for densities, distribution functions, quantiles, survival probabilities, and tail risks. The framework is conditional by construction; unconditional models arise as a special case when covariates are omitted. The package supports both Chinese restaurant process (CRP) and stick-breaking (SB) representations of the Dirichlet process, multiple kernel families for bulk modeling, and an extension to causal inference via treatment-specific conditional distributions that enable inference on distributional and tail-sensitive causal estimands.

Keywords: Bayesian nonparametrics; Dirichlet process mixture; generalized Pareto distribution; extreme value theory; posterior prediction; causal inference; quantile treatment effects; R; `nimble`.

1. Introduction

Heavy-tailed outcomes arise in many empirical settings where rare events dominate risk assessment and decision-making. Examples include environmental and climate applications, where extremes of precipitation or temperature shape planning and infrastructure design (Coles

2001a; Katz, Parlange, and Naveau 2002; Davison and Smith 1990), insurance and finance where catastrophic losses drive aggregate exposure (Embrechts, Klüppelberg, and Mikosch 1997a), and biomedical and public health research where infrequent but severe outcomes may outweigh average effects (Austin and Schuster 2016). Across these domains, analysts frequently require inference for extreme quantiles, tail probabilities, and other functionals of the full outcome distribution rather than only its mean.

Standard parametric models are often inadequate for such tasks because they impose restrictive assumptions on tail behavior. Misspecification in the tail can induce substantial bias in estimated exceedance probabilities and extreme quantiles even when the model appears adequate in the bulk (Coles 2001a; Davison and Smith 1990). Fully nonparametric approaches offer substantial flexibility in regions where data are abundant, but they typically provide limited guidance for extrapolation in the tail where observations are sparse and uncertainty must be carefully controlled (Ghosal and van der Vaart 2017).

Extreme value theory (EVT) provides a mathematically principled framework for modeling rare events. Classical results due to Balkema and de Haan (1974) and Pickands (1975a) show that, under mild regularity conditions, threshold exceedances converge in distribution to a generalized Pareto distribution (GPD). This asymptotic justification underlies the widespread use of peaks-over-threshold modeling in applied work (Davison and Smith 1990; Coles 2001a; de Haan and Ferreira 2007). In R, EVT methodology is supported by packages such as `evir` (Pfaff 2018), `evd` (Stephenson 2002), `ismev` (Heffernan and Stephenson 2018), and `extRemes` (Gilleland and Katz 2016), among others. However, EVT focuses on extremes and does not specify a model for the bulk distribution, which is often handled heuristically or through ad hoc parametric assumptions.

Bayesian nonparametric methods, particularly Dirichlet process mixtures (DPMs), provide a complementary approach for flexible bulk modeling. The Dirichlet process introduced by Ferguson (1973a) induces a random discrete mixing distribution and has become foundational in Bayesian nonparametrics. Dirichlet process mixtures are widely used for density estimation and clustering (Lo 1984a; Escobar and West 1995; Neal 2000a), and are implemented in R through packages such as `DPpackage` (Jara, Hanson, Quintana, Müller, and Rosner 2011), `dirichletprocess` (Ross and Markwick 2023), and `BNPdensity` (Arbel and Canale 2021). Probabilistic programming tools, notably `nimble` (de Valpine, Turek, Paciorek, Anderson-Bergman, Temple Lang, and Bodik 2017), further enable custom hierarchical Bayesian models and bespoke MCMC algorithms within R. Despite these advantages, standard DPMs inherit their tail behavior from the chosen kernel family, which is typically selected for bulk fit rather than for asymptotic tail properties, and this can yield unstable inference for extreme quantiles.

These considerations motivate semiparametric models that combine flexible bulk modeling with EVT-based tails. Several Bayesian approaches combine nonparametric or flexible bulk models with parametric tails to improve tail inference while retaining bulk flexibility (do Nascimento, Gamerman, and Lopes 2012; Fúquene-Patiño, Müller, and Rosner 2015; Cabras and Castellanos 2011). In particular, Fúquene-Patiño *et al.* (2015) propose a Dirichlet process mixture of gamma densities for the bulk with a generalized Pareto tail for exceedances, demonstrating stable posterior inference for high quantiles. However, available implementations are often limited in scope, focus on unconditional models, or do not provide a unified and extensible software suite for conditional modeling, posterior prediction, and causal extensions.

Heavy tails also create challenges for causal inference, where treatment effects may differ across the outcome distribution and where tail behavior can be directly policy-relevant. Distributional causal estimands such as quantile treatment effects (Firpo, Fortin, and Lemieux 2009; Chernozhukov, Fernández-Val, and Melly 2013) and tail-risk contrasts can be more informative than average treatment effects when extremes dominate costs or harms. In R, causal and distributional methods are supported by software such as `quantreg` (Koenker 2005), `grf` (Athey, Tibshirani, and Wager 2019), and related packages for propensity score weighting and modern causal modeling. Nevertheless, many causal workflows rely on outcome models without explicit EVT regularization, which can limit reliability for extreme outcomes.

This article introduces `DPmixGPD`, an R package implementing a Bayesian semiparametric framework that combines conditional Dirichlet process mixtures for the bulk distribution with covariate-dependent generalized Pareto tails for exceedances, and provides a causal extension via treatment-specific conditional distributions. The framework is conditional by construction; unconditional models arise as a special case when covariates are omitted. Posterior inference jointly accounts for uncertainty in the bulk mixture, tail parameters, and the threshold separating bulk and tail regimes, enabling stable posterior predictive inference for densities, distribution functions, quantiles, survival probabilities, and tail risks.

The remainder of the article is organized as follows. Section **??** reviews background on conditional DPMs, CRP and stick-breaking representations, and EVT-based tail modeling. Section 3 presents the conditional bulk–tail model with a complete likelihood and prior specification. Section 4 details posterior computation and posterior predictive inference under both CRP and stick-breaking implementations. Section 5 develops the causal extension and the corresponding posterior predictive estimands. Section 6 describes the software architecture and user-facing interface of `DPmixGPD`. Simulation studies and real-data applications illustrate performance and practical utility.

2. Background and Related Work

The methodology implemented in `DPmixGPD` combines three foundational areas of statistical theory: Bayesian nonparametric mixture modeling for flexible conditional distributions, extreme value theory for tail modeling, and distributional causal inference for principled comparisons across treatment regimes. While each area is well developed individually, they are seldom integrated into a unified inferential and software framework. This section provides a review of the mathematical foundations for each component and describes their integration in `DPmixGPD`.

2.1. Conditional Dirichlet Process Mixture Models

Let $Y \in \mathbb{R}$ denote a univariate outcome, and let $\mathbf{X} = \mathbf{x} \in \mathbb{R}^p$ denote observed covariates. The goal is to estimate the conditional distribution function $F(y \mid \mathbf{x})$ and its associated conditional density $p(y \mid \mathbf{x})$ without restrictive parametric assumptions. Bayesian nonparametric approaches achieve this via flexible mixture models whose complexity adapts to the data.

The conditional Dirichlet process mixture (DPM) model expresses the conditional density as

an infinite mixture of kernel densities:

$$F_{\mathrm{DPM}}(y \mid \mathbf{x}) = \sum_{j=1}^{\infty} w_j \, F_j(y \mid \mathbf{x}),$$

where each component distribution $F_j(y \mid \mathbf{x})$ is associated with a kernel $f(y \mid \theta_j(\mathbf{x}))$ and component-specific parameters θ_j, while the weights $\{w_j\}$ arise from a Dirichlet process prior (Ferguson 1973b; Lo 1984b). The covariate dependence enters through $\theta_j(\mathbf{x})$, often modeled using a regression structure with link functions to enforce parameter constraints.

This construction supports localized adaptation of the distributional shape across the covariate space (Müller and Rosner 1996; Gelfand, Kottas, and MacEachern 2005). In `DPmixGPD`, F_{DPM} defines the "bulk" of the outcome distribution, capturing features such as multimodality and heteroscedasticity in a covariate-dependent manner.

2.2. CRP and Stick-Breaking Representations

Two alternative representations of the Dirichlet process are particularly useful for computation: the Chinese restaurant process (CRP) and the stick-breaking (SB) construction. Both are implemented in `DPmixGPD` to provide algorithmic flexibility.

In the CRP representation, the random mixing measure is integrated out, and inference proceeds by sampling cluster assignments z_i for each observation. The CRP induces the predictive rule:

$$\Pr(z_i = j \mid z_1, \ldots, z_{i-1}) = \begin{cases} \frac{n_j}{i-1+\alpha}, & \text{if cluster } j \text{ exists}, \\ \frac{\alpha}{i-1+\alpha}, & \text{for a new cluster}, \end{cases}$$

where n_j is the number of prior observations in cluster j. Given $z_i = j$, the likelihood contribution is $f(y_i \mid \theta_j(\mathbf{x}_i))$. This marginal approach often improves mixing in MCMC sampling (Neal 2000b).

Alternatively, the stick-breaking construction represents the random measure explicitly:

$$G = \sum_{j=1}^{\infty} w_j \delta_{\theta_j}, \quad w_j = v_j \prod_{\ell<j} (1 - v_\ell), \quad v_j \sim \mathrm{Beta}(1, \alpha), \quad \theta_j \sim G_0,$$

as introduced by Sethuraman (1994). This form facilitates blocked Gibbs samplers and variational methods (Ishwaran and James 2001). In `DPmixGPD`, either representation may be used to define the conditional distribution $F_{\mathrm{DPM}}(y \mid \mathbf{x})$, depending on the computational context.

2.3. Extreme Value Theory and Generalized Pareto Tails

While DPM models offer flexibility in the bulk, their tail behavior is determined by the kernel family and is not explicitly modeled. Extreme value theory (EVT) provides a principled framework for modeling the distribution of excesses above a high threshold.

Given a covariate value $\mathbf{x}$, define a threshold $u(\mathbf{x})$, and consider the excess $E = Y - u(\mathbf{x})$ given $Y > u(\mathbf{x})$. EVT shows that, under mild conditions, the conditional distribution of E converges to the generalized Pareto distribution (GPD) (Pickands 1975b; Embrechts, Klüppelberg, and Mikosch 1997b):

$$F_{\mathrm{GPD}}(e \mid \sigma(\mathbf{x}), \xi(\mathbf{x})) = 1 - \left(1 + \xi(\mathbf{x}) \frac{e}{\sigma(\mathbf{x})}\right)^{-1/\xi(\mathbf{x})}, \quad e > 0,$$

with associated density

$$g(e \mid \sigma(\mathbf{x}), \xi(\mathbf{x})) = \frac{1}{\sigma(\mathbf{x})} \left(1 + \xi(\mathbf{x}) \frac{e}{\sigma(\mathbf{x})}\right)^{-1/\xi(\mathbf{x})-1},$$

defined on the support where $1 + \xi(\mathbf{x})e/\sigma(\mathbf{x}) > 0$. The scale parameter $\sigma(\mathbf{x})$ controls dispersion, while the shape parameter $\xi(\mathbf{x})$ governs tail heaviness.

`DPmixGPD` models both $\sigma(\mathbf{x})$ and $\xi(\mathbf{x})$ as functions of covariates using link-regression structures, following approaches similar to Coles (2001b), allowing tail behavior to vary across the covariate space.

2.4. Conditional Bulk–Tail Splicing

To obtain a unified model, `DPmixGPD` employs a conditional bulk–tail splicing strategy. The overall conditional distribution is defined as:

$$F(y \mid \mathbf{x}) = \begin{cases} F_{\mathrm{DPM}}(y \mid \mathbf{x}), & y \leq u(\mathbf{x}), \\ F_{\mathrm{DPM}}(u(\mathbf{x}) \mid \mathbf{x}) + [1 - F_{\mathrm{DPM}}(u(\mathbf{x}) \mid \mathbf{x})] F_{\mathrm{GPD}}(y - u(\mathbf{x}) \mid \sigma(\mathbf{x}), \xi(\mathbf{x})), & y > u(\mathbf{x}). \end{cases}$$

This formulation guarantees continuity at the threshold and preserves the total probability mass. By integrating nonparametric modeling in the center and parametric modeling in the tail, it enables coherent estimation of the full distribution. Similar splicing constructions have been proposed in both frequentist and Bayesian contexts (Tancredi and Anderson 2006; Carreau and Bengio 2009; Naveau, Huser, Ribereau, and Hannart 2016).

2.5. Causal Structure and Modeling Framework

Causal inference is formulated using the potential outcomes framework (Rubin 1974; Imbens and Rubin 2015). Let $Y^{(1)}$ and $Y^{(0)}$ denote the potential outcomes under treatment and control, respectively, and let $A \in \{0, 1\}$ be the observed treatment assignment. We assume a set of baseline covariates $\mathbf{X} = \mathbf{x}$ is available for each unit. Identification of the conditional distributions of potential outcomes relies on the following standard assumptions:

1. **Consistency**: $Y = Y^{(A)}$.
2. **No interference**: One unit's potential outcomes do not depend on the treatment assignment of other units.
3. **Conditional exchangeability**: $\{Y^{(0)}, Y^{(1)}\} \perp A \mid \mathbf{X}$.
4. **Positivity**: $0 < \Pr(A = 1 \mid \mathbf{X} = \mathbf{x}) < 1$ for all $\mathbf{x}$ in the covariate support.

Under these assumptions, the conditional distributions $F^{(1)}(y \mid \mathbf{x})$ and $F^{(0)}(y \mid \mathbf{x})$ are identifiable from observed data. In `DPmixGPD`, estimation proceeds by modeling these two potential outcome distributions separately using the same spliced DPM–GPD structure. Covariate adjustment is achieved via propensity score methods (Rosenbaum and Rubin 1983), where the propensity score $\pi(\mathbf{x}) = \Pr(A = 1 \mid \mathbf{X} = \mathbf{x})$ is estimated and used to adjust for confounding.

This modeling scheme supports flexible, covariate-dependent estimation of both marginal and extreme behavior under each treatment regime, and provides a coherent basis for distributional causal inference.

3. Model Development and Posterior Summary Calculation

4. Posterior Computation and Posterior Predictive Inference

5. Causal Extension

6. Software

7. Simulation Studies

8. Real-Data Applications

Appendix: Example Code Skeleton (not executed)

The following R code chunks illustrate the intended workflow and are included for documentation purposes only. They are not executed in this manuscript.

```
# library(DPmixGPD)

# fit <- dpmixgpd(
#   y = y,
#   X = X,
#   backend = "sb",         # or "crp"
#   kernel  = "gamma",      # e.g., "normal", "lognormal", "invgauss", "laplace", "cauchy",
#   gpd     = TRUE,
#   mcmc    = list(niter = 5000, nburnin = 1000, thin = 5, nchains = 2, seed = c(1,2))
# )
# summary(fit)
# plot(fit)
```

```
# pred <- predict(
#   fit,
#   newdata = Xnew,
#   type = c("density","cdf","quantile","survival"),
#   probs = c(0.5, 0.9, 0.99)
# )
# pred
```

References

Arbel J, Canale A (2021). "BNPdensity: Bayesian nonparametric density estimation." *Journal of Statistical Software*.

Athey S, Tibshirani J, Wager S (2019). "Generalized random forests." *The Annals of Statistics*, **47**(2), 1148–1178.

Austin PC, Schuster T (2016). "Estimating the effect of treatment on rare outcomes." *Statistics in Medicine*, **35**(24), 4246–4260.

Balkema AA, de Haan L (1974). "Residual life time at great age." *The Annals of Probability*, **2**(5), 792–804.

Cabras S, Castellanos ME (2011). "A Bayesian hierarchical model for extreme rainfall events." *Environmetrics*, **22**(1), 13–28.

Carreau J, Bengio Y (2009). "A hybrid Pareto model for asymmetric fat-tailed data: The univariate case." *Extremes*, **12**(1), 29–44.

Chernozhukov V, Fernández-Val I, Melly B (2013). "Inference on counterfactual distributions." *Econometrica*, **81**(6), 2205–2268.

Coles S (2001a). *An Introduction to Statistical Modeling of Extreme Values*. Springer.

Coles S (2001b). *An Introduction to Statistical Modeling of Extreme Values*. Springer.

Davison AC, Smith RL (1990). "Models for exceedances over high thresholds." *Journal of the Royal Statistical Society: Series B*, **52**(3), 393–442.

de Haan L, Ferreira A (2007). *Extreme Value Theory: An Introduction*. Springer.

de Valpine P, Turek D, Paciorek C, Anderson-Bergman C, Temple Lang D, Bodik R (2017). "Programming with models in NIMBLE." *Journal of Computational and Graphical Statistics*, **26**(2), 403–413.

do Nascimento FF, Gamerman D, Lopes HF (2012). "A Bayesian semiparametric model for extreme rainfall events." *Journal of Agricultural, Biological, and Environmental Statistics*, **17**(2), 253–273.

Embrechts P, Klüppelberg C, Mikosch T (1997a). *Modelling Extremal Events for Insurance and Finance*. Springer.

Embrechts P, Klüppelberg C, Mikosch T (1997b). *Modelling Extremal Events for Insurance and Finance*. Springer.

Escobar MD, West M (1995). "Bayesian density estimation and inference using mixtures." *Journal of the American Statistical Association*, **90**(430), 577–588.

Ferguson TS (1973a). "A Bayesian analysis of some nonparametric problems." *The Annals of Statistics*, **1**(2), 209–230.

Ferguson TS (1973b). "A Bayesian analysis of some nonparametric problems." *The Annals of Statistics*, **1**(2), 209–230.

Firpo S, Fortin NM, Lemieux T (2009). "Unconditional quantile regressions." *Econometrica*, **77**(3), 953–973.

Fúquene-Patiño J, Müller P, Rosner G (2015). "A Bayesian semiparametric approach for modeling heavy-tailed data." *Biometrics*, **71**(2), 337–346.

Gelfand AE, Kottas A, MacEachern SN (2005). "Bayesian nonparametric spatial modeling with Dirichlet process mixing." *Journal of the American Statistical Association*, **100**(471), 1021–1035.

Ghosal S, van der Vaart A (2017). *Fundamentals of Nonparametric Bayesian Inference*. Cambridge University Press.

Gilleland E, Katz RW (2016). "extRemes 2.0." *Journal of Statistical Software*, **72**(8), 1–39.

Heffernan J, Stephenson A (2018). *ismev: An Introduction to Statistical Modeling of Extreme Values*. R package.

Imbens GW, Rubin DB (2015). *Causal Inference for Statistics, Social, and Biomedical Sciences: An Introduction*. Cambridge University Press.

Ishwaran H, James LF (2001). "Gibbs sampling methods for stick-breaking priors." *Journal of the American Statistical Association*, **96**(453), 161–173.

Jara A, Hanson T, Quintana F, Müller P, Rosner G (2011). "DPpackage: Bayesian nonparametric modeling in R." *Journal of Statistical Software*, **40**(5), 1–30.

Katz RW, Parlange MB, Naveau P (2002). "Statistics of extremes in hydrology." *Advances in Water Resources*, **25**(8–12), 1287–1304.

Koenker R (2005). *Quantile Regression*. Cambridge University Press.

Lo AY (1984a). "On a class of Bayesian nonparametric estimates." *The Annals of Statistics*, **12**(1), 351–357.

Lo AY (1984b). "On a class of Bayesian nonparametric estimates: I. Density estimates." *The Annals of Statistics*, **12**(1), 351–357.

Müller P, Rosner GL (1996). "Bayesian hierarchical modeling for clustered data with application to growth curves." *Biometrics*, **52**(4), 963–973.

Naveau P, Huser R, Ribereau P, Hannart A (2016). "Modeling heavy rainfall using a Bayesian peaks-over-threshold approach." *Weather and Climate Extremes*, **13**, 1–10.

Neal RM (2000a). "Markov chain sampling methods for Dirichlet process mixture models." *Journal of Computational and Graphical Statistics*, **9**(2), 249–265.

Neal RM (2000b). "Markov chain sampling methods for Dirichlet process mixture models." *Journal of Computational and Graphical Statistics*, **9**(2), 249–265.

Pfaff B (2018). *evir: Extreme Values in R*. R package.

Pickands J (1975a). "Statistical inference using extreme order statistics." *The Annals of Statistics*, **3**(1), 119–131.

Pickands J (1975b). "Statistical inference using extreme order statistics." *The Annals of Statistics*, **3**(1), 119–131.

Rosenbaum PR, Rubin DB (1983). "The central role of the propensity score in observational studies for causal effects." *Biometrika*, **70**(1), 41–55.

Ross GJ, Markwick D (2023). "dirichletprocess: Build Dirichlet process objects in R." *Journal of Statistical Software*.

Rubin DB (1974). "Estimating causal effects of treatments in randomized and nonrandomized studies." *Journal of Educational Psychology*, **66**(5), 688–701.

Sethuraman J (1994). "A constructive definition of Dirichlet priors." *Statistica Sinica*, **4**, 639–650.

Stephenson A (2002). "evd: Extreme value distributions." *R News*, **2**(2), 31–32.

Tancredi A, Anderson CW (2006). "A hierarchical Bayesian spliced model for the assessment of the probability of extreme wind speeds." *Environmetrics*, **17**(3), 209–229.

Affiliation:

Firstname Lastname
Affiliation
Address, Country
E-mail: name@address
URL: http://link/to/webpage/

Journal of Statistical Software http://www.jstatsoft.org/
published by the Foundation for Open Access Statistics http://www.foastat.org/
MMMMMM YYYY, Volume VV, Issue II *Submitted:* yyyy-mm-dd
doi:10.18637/jss.v000.i00 *Accepted:* yyyy-mm-dd